

Seasonality Clustering

A Hierarchical Agglomerative Approach

Intern Project – Summer 2025

08/13/2025

Adam Mahmoud



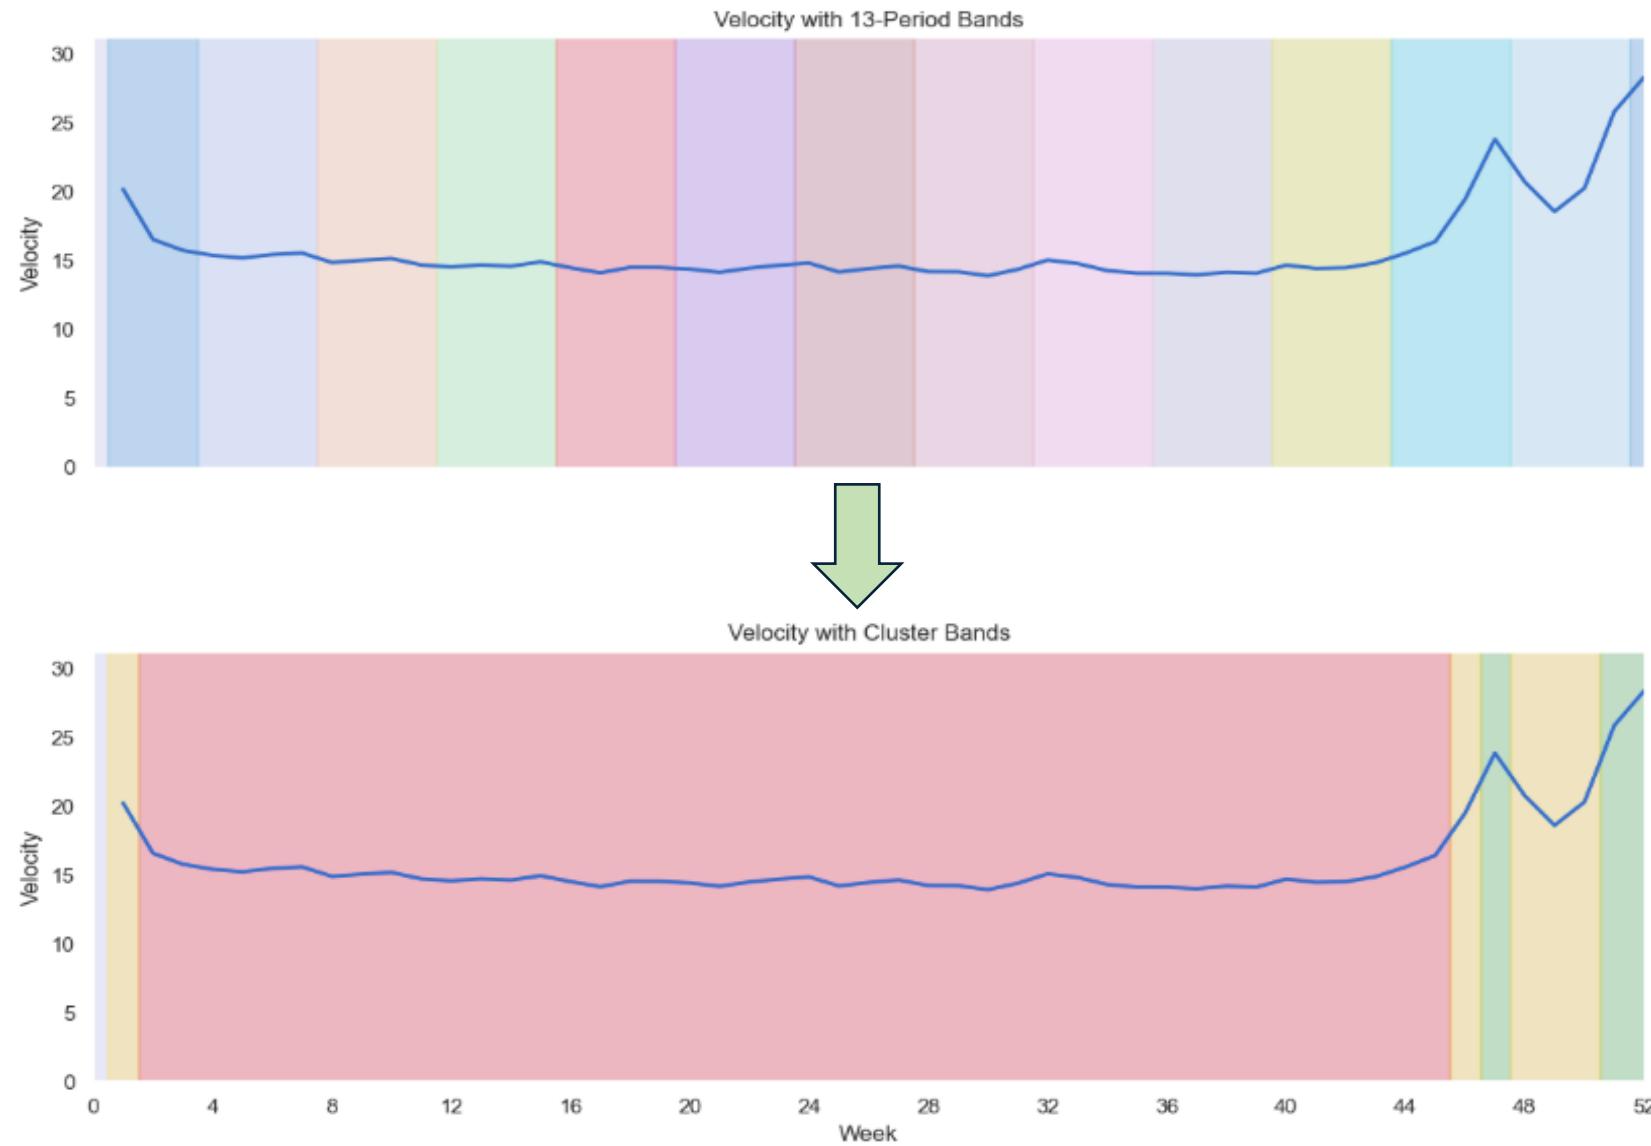
Background & Context

- Current pricing regressions incorporate seasonality through 13 uniform quad periods per year, roughly aligned to holidays.
 - Quad boundaries may misalign with true demand patterns → risk of capturing noise instead of seasonality
- Clustering weeks into fewer data-driven seasons can:
 - Better reflect actual demand trends
 - Reduce model complexity & collinearity
 - Improve interpretability
- Bases season definitions on client-specific data

Executive Summary

- **Problem Statement:** How can we replace fixed quad-period seasonality with data-driven seasonal clusters to reduce pricing regression complexity and capture true demand variation.
- **Hierarchical Agglomerative Clustering**
 - Base Dollar Velocity
 - Base Dollar Velocity & Time
- **Evaluation Metrics**
 - Silhouette Score, Calinski-Harabasz (CH) Index, Davies-Bouldin (DB) Index
 - Dendrogram
 - Silhouette Plot
 - *Velocity & Time Clustering only: Silhouette Score vs. Alpha plots*
- **Regressions**
 - Compare Demand Indices and regression model metrics between pricing regressions that:
 - do not use period
 - use standard quad-periods
 - use clustered periods

Visual Overview: Quad Periods vs. Clustered Seasons – Martinelli's Data



Data & Attributes

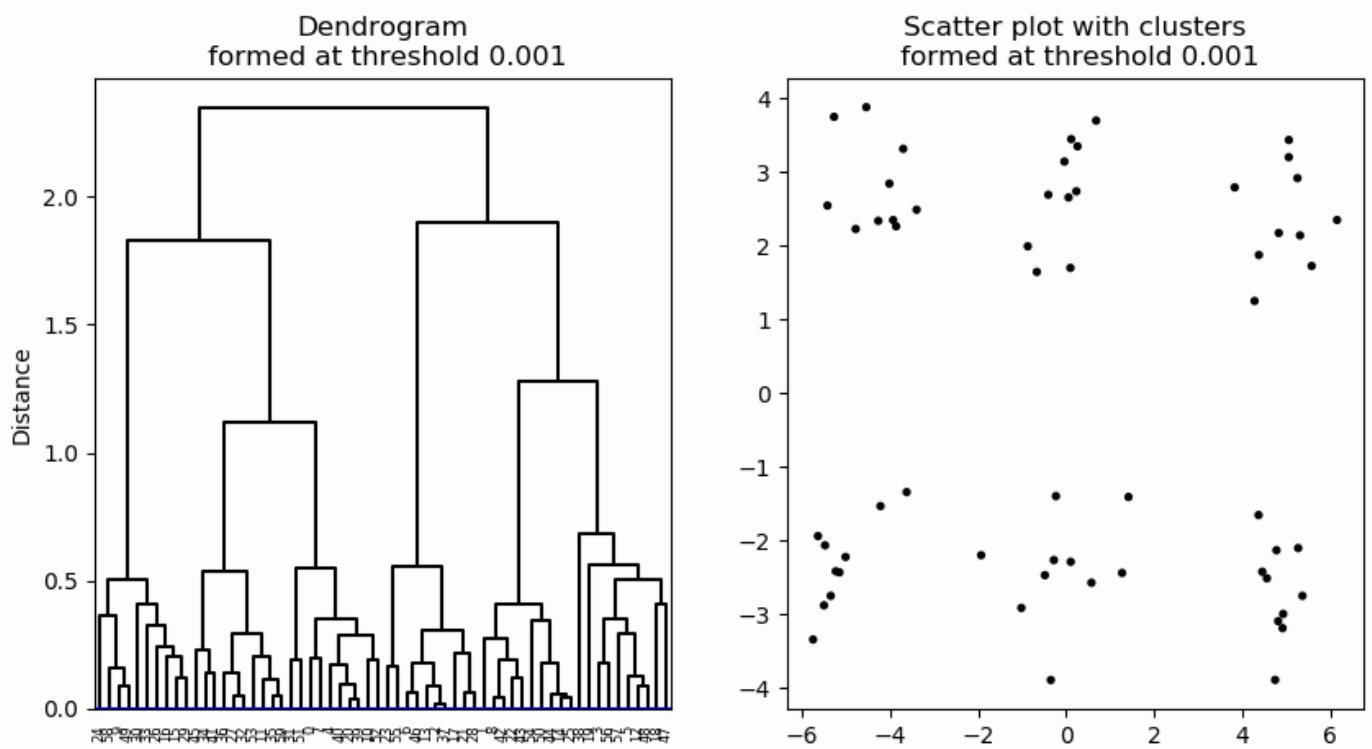
- **Datasets for building clustering pipeline:**
 - Chomps (low seasonality)
 - TruFru (medium seasonality)
 - Martinelli's (high seasonality)
- **Attributes:**
 - Account: Total US Food
 - Base Dollar Velocity: Base Dollars / Stores Selling / Weeks in Distribution
 - Velocity is aggregated over multiple years
 - Note: Database baselines are used

Methodology

- Hierarchical agglomerative clustering
- Clustering on velocity
- Clustering on velocity & time
 - Alpha optimization
- Regression Validation

Hierarchical Agglomerative Clustering

- **Unsupervised machine learning:** no input-output pairs, no period labels on week numbers
- **Agglomerative:** every data point in its own cluster → merge *similar* pairs of clusters until 1 is left
- **Linkage Criterion:** methods for deciding the order of cluster combinations
 - **Single**, complete, average, weighted, centroid, median, ward



Cluster Evaluation Metrics

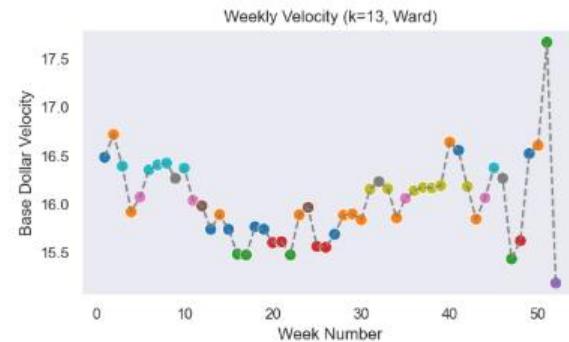
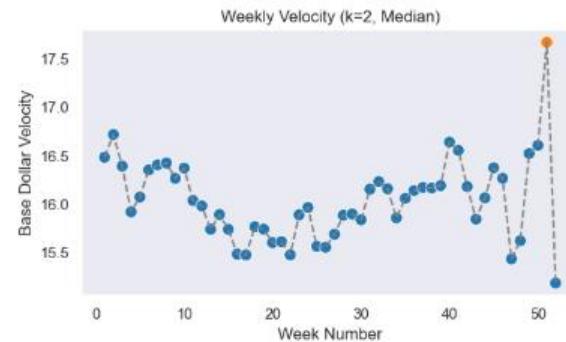
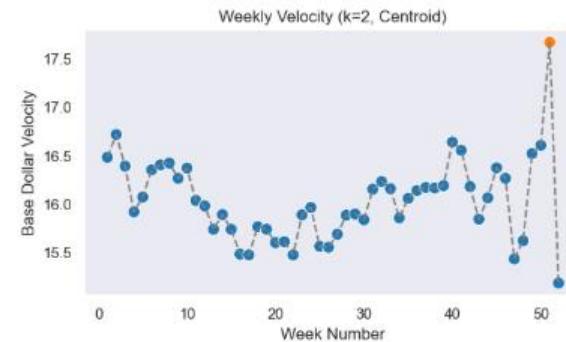
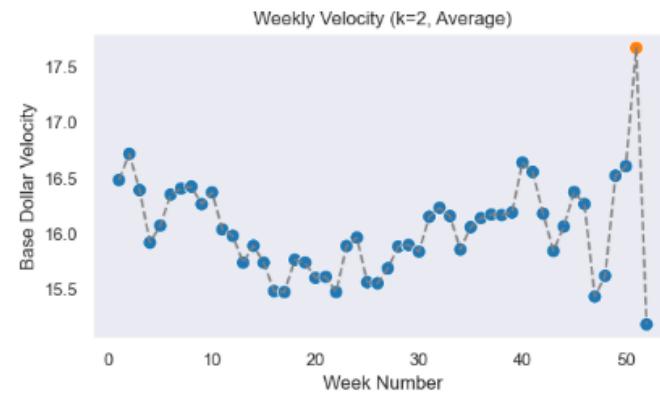
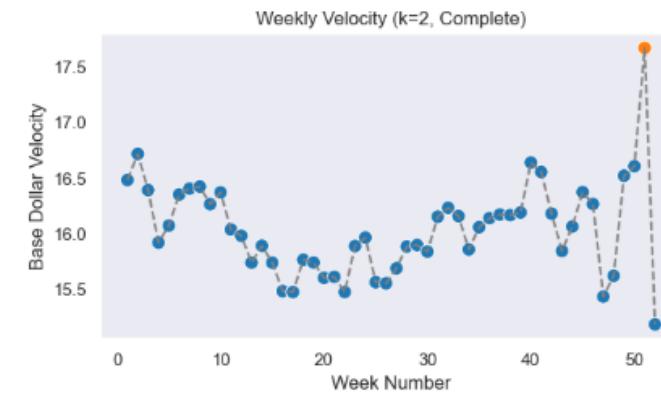
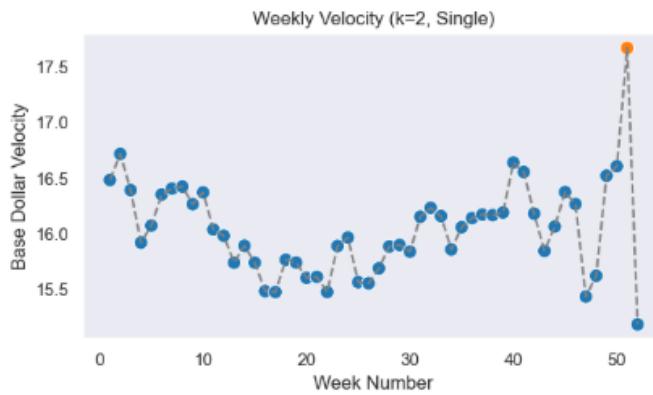
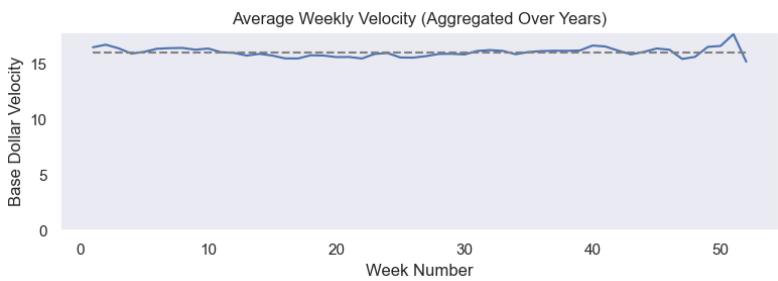
- Dendrogram: tree diagram showing how clusters merge step-by-step; merge height reflects the distance between joined clusters.
- **Silhouette Score:** measures how similar a point is to its own cluster compared to other clusters. Ranges from -1 (poor fit) to 1 (well separated); higher is better.
- CH Index: ratio of between-cluster variance to within-cluster variance, adjusted for number of clusters. Higher values indicate better defined clusters.
- DB Index: measures average similarity between each cluster and its most similar other cluster; lower values indicate tighter, more distinct clusters.

Clustering on Velocity

Case Study: Chomps

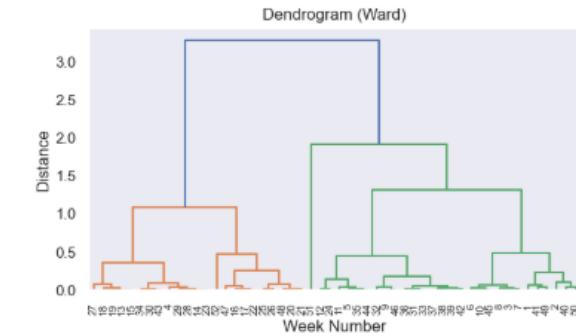
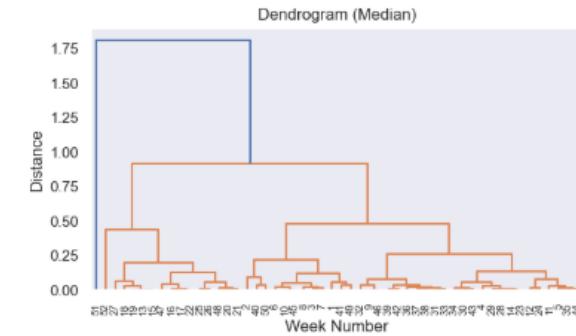
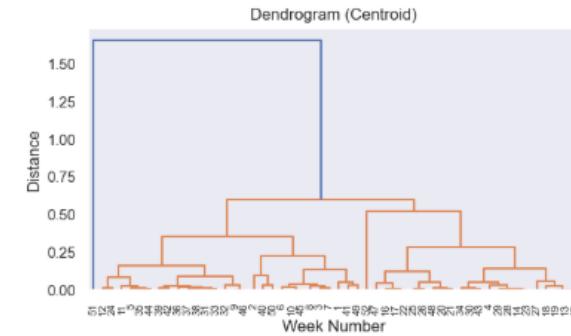
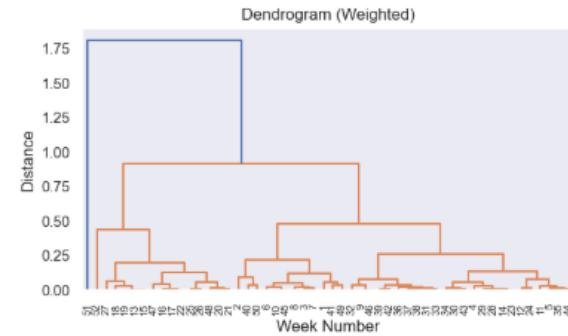
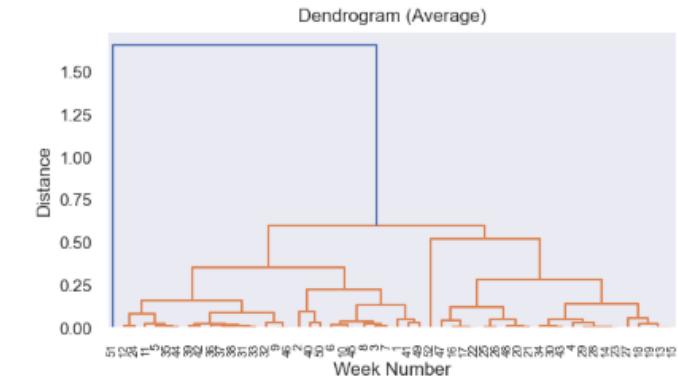
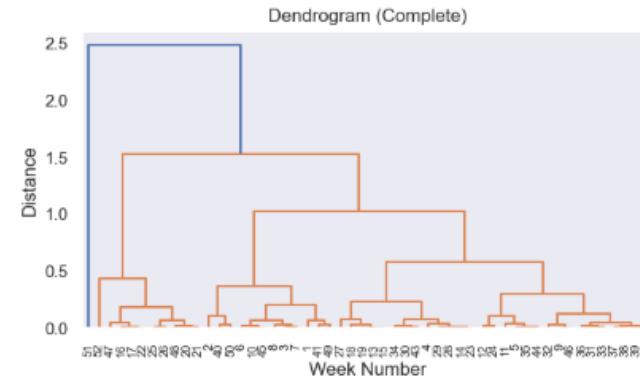
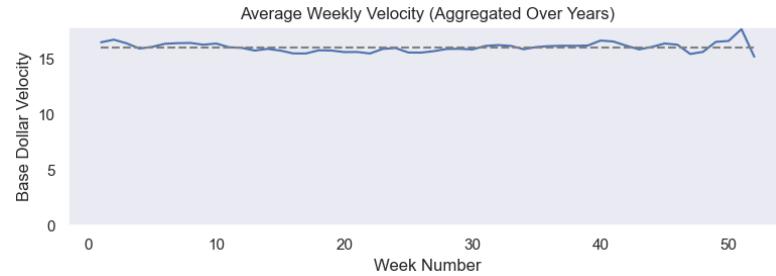
Takeaway: Narrow focus to Single Linkage & Silhouette Score

Chomps – Velocity Clustering



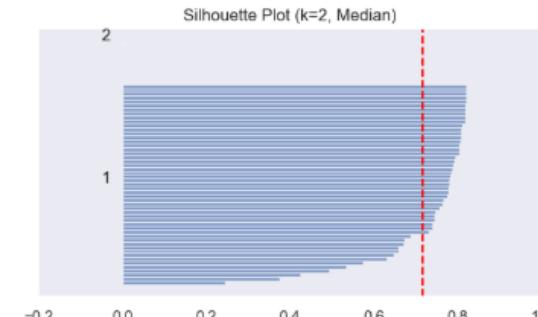
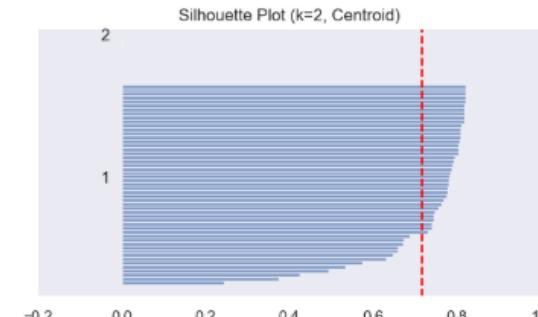
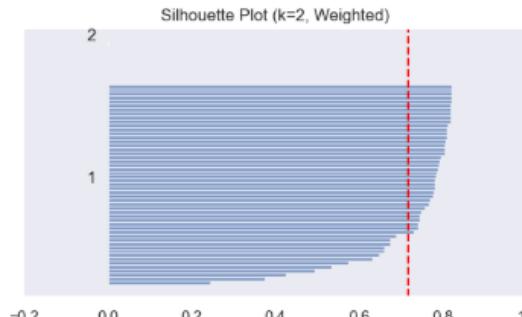
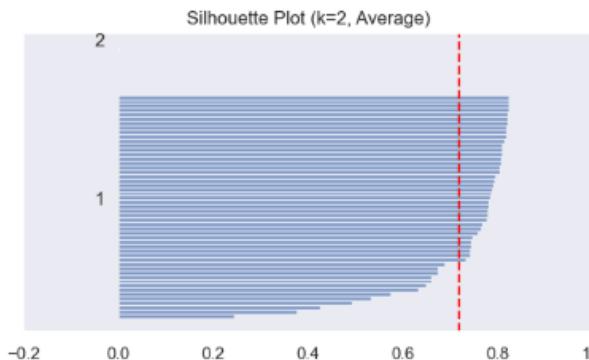
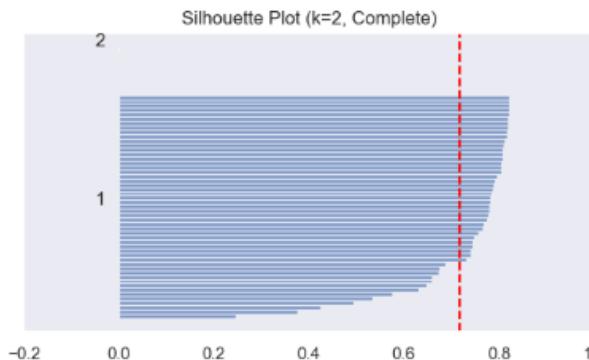
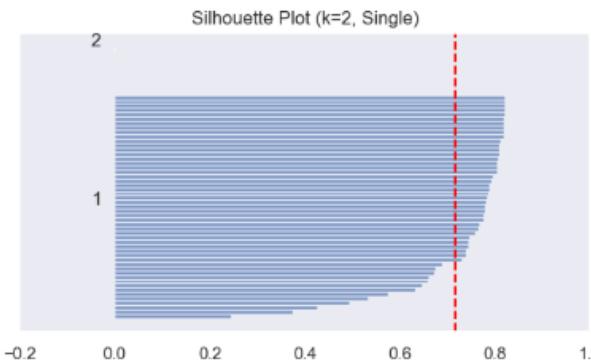
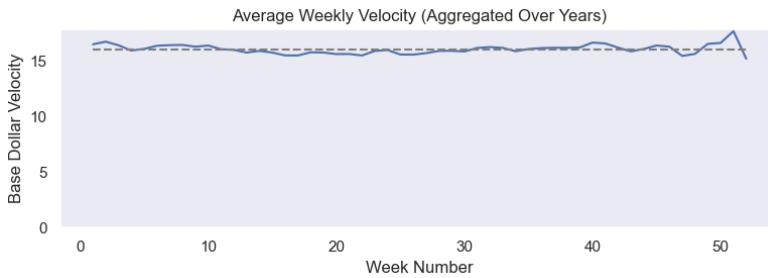
Note: $k_{\max} = 13$

Chomps – Dendrograms



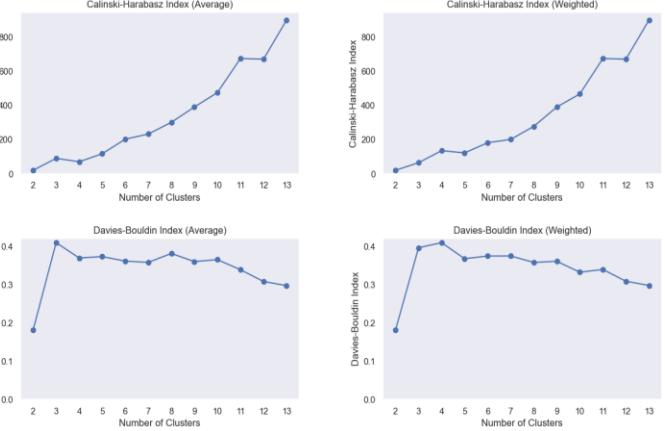
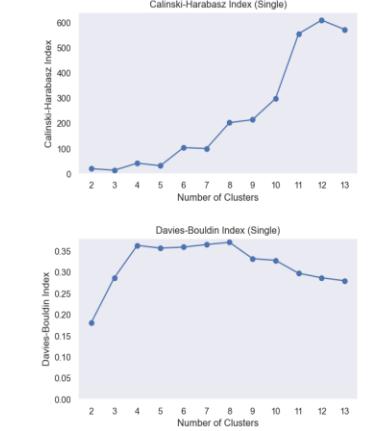
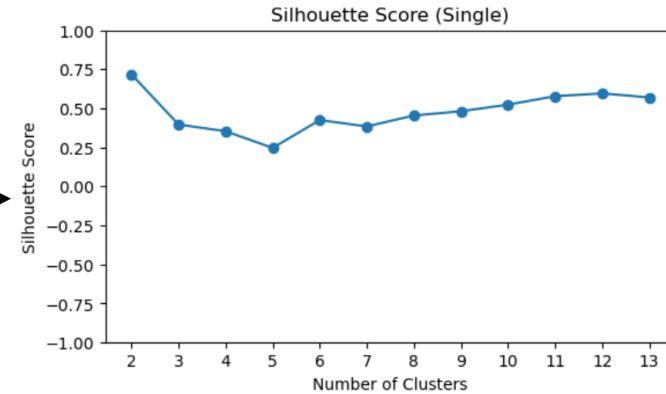
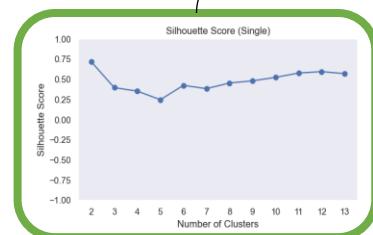
Note: k_{max} = 13

Chomps – Silhouette Plots



Note: k_max = 13

Chomps – Silhouette Score, CH Index, DB Index by Number of Clusters



Note: k_max = 13

Clustering on Velocity & Time

Case Study: TruFru

Takeaway: Incorporating time in a custom distance function typically worsens clustering

Custom Distance Functions Tested

- ❖ **Linear:** Combines normalized week difference and velocity difference with a weighted average
- **Squared Velocity:** Same as linear, but velocity difference is squared to emphasize larger gaps
- **Exponential Decay on Time:** Uses exponential decay for time difference, making nearby weeks much closer
- **Cosine Bump Distance:** Gives extra closeness to weeks within 4 using a cosine curve, then switches to linear growth; piecewise function

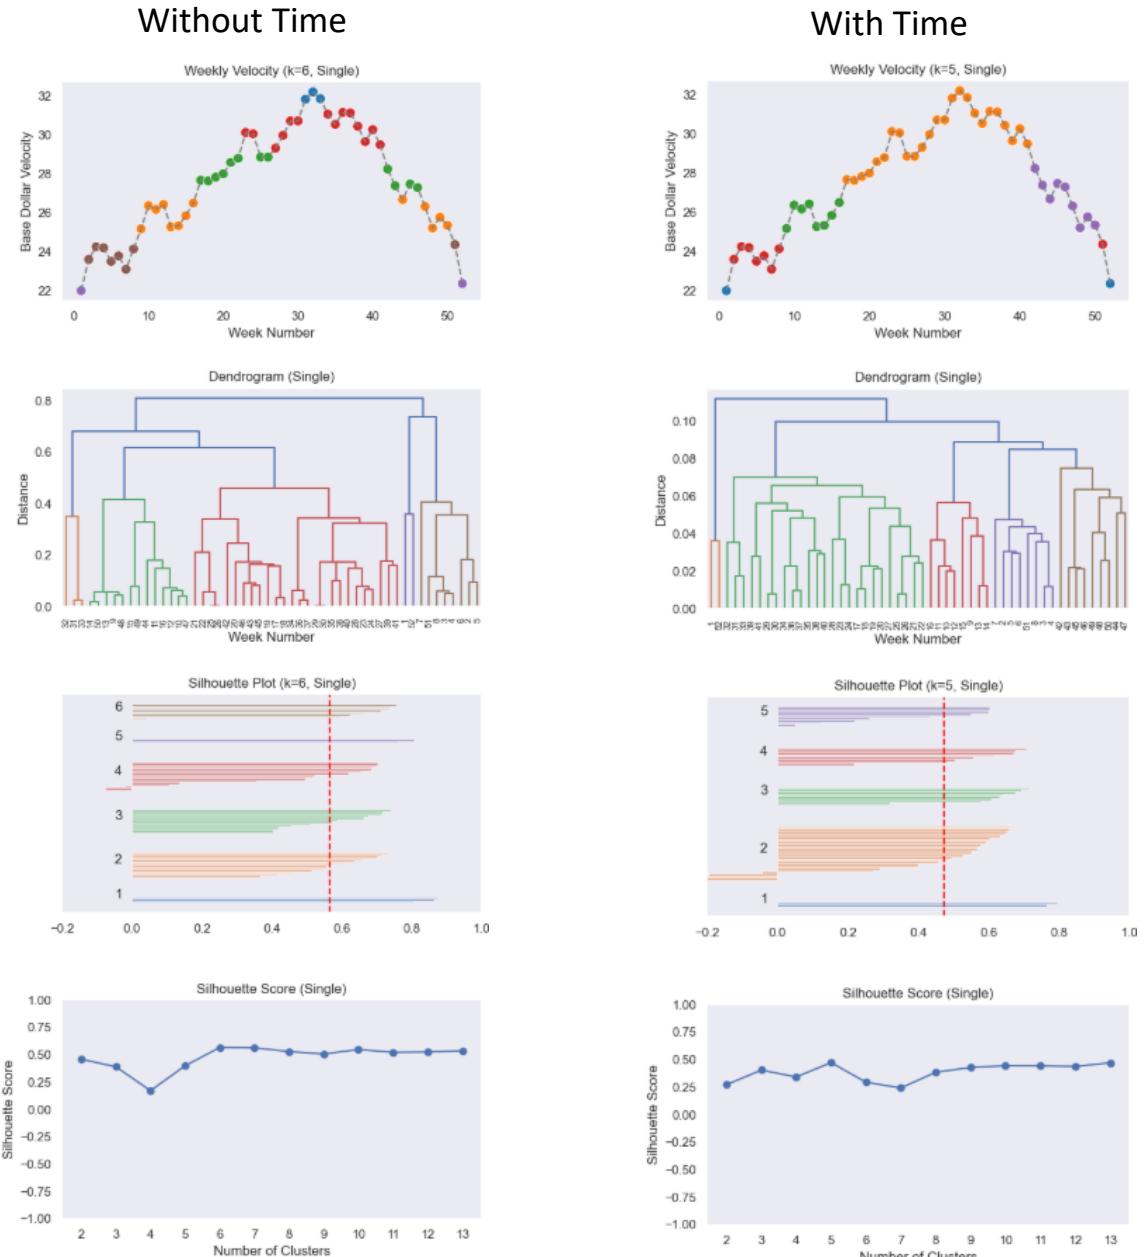
TruFru – Clustering With & Without Time

- **Linear Distance Function:**

$$d = \alpha \cdot \frac{\text{week_diff}}{26} + (1 - \alpha) \cdot \frac{|v_1 - v_2|}{v_{\max} - v_{\min}}$$

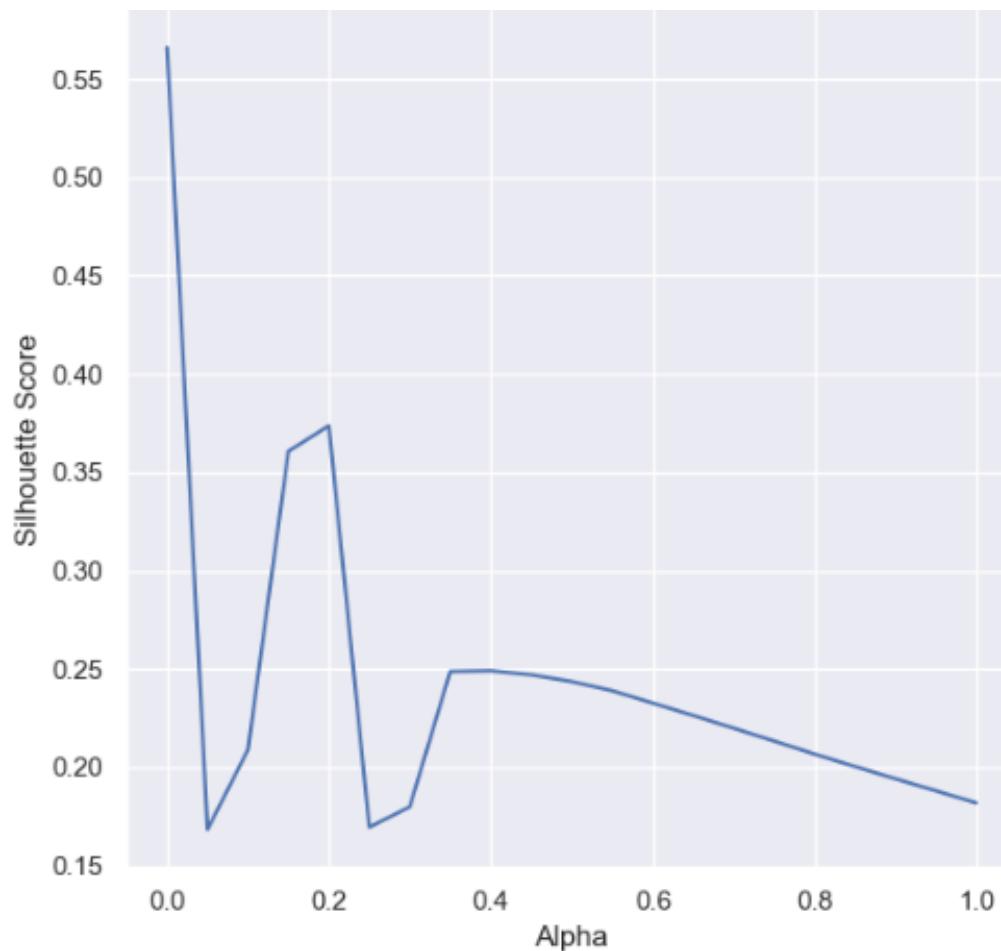
where $\alpha = 0.2$

- Generally lower silhouette scores
- Clustering with time led to an optimized silhouette score at a lower number of clusters



TruFru – Increasing Alpha (putting more weight to the time component in calculating distance) generally worsens average silhouette score

- Linear distance function
- Single linkage
- Number of Clusters = 6

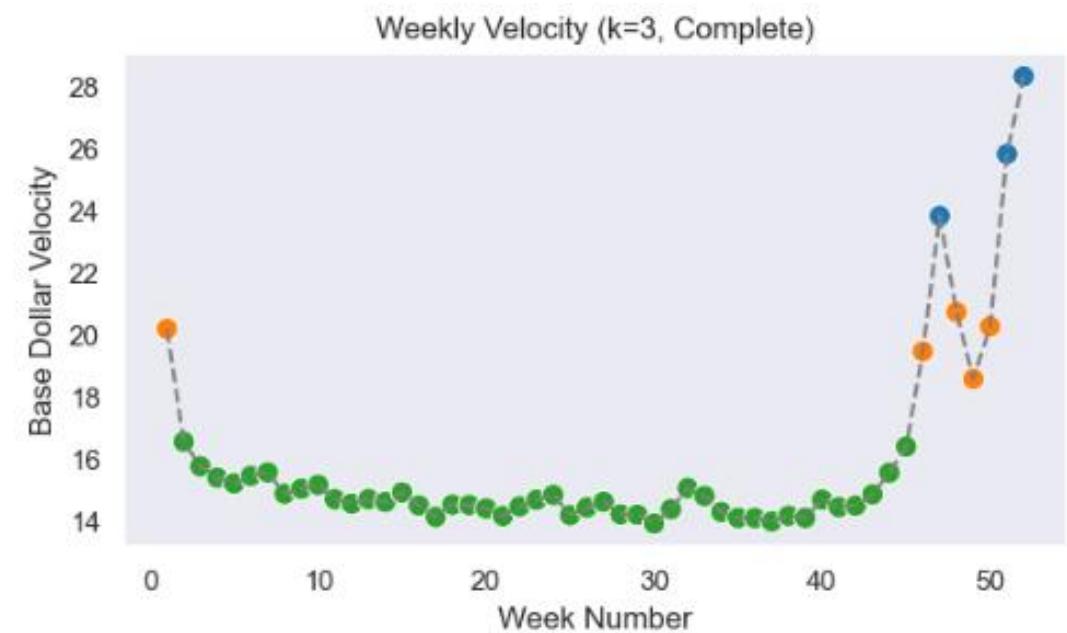
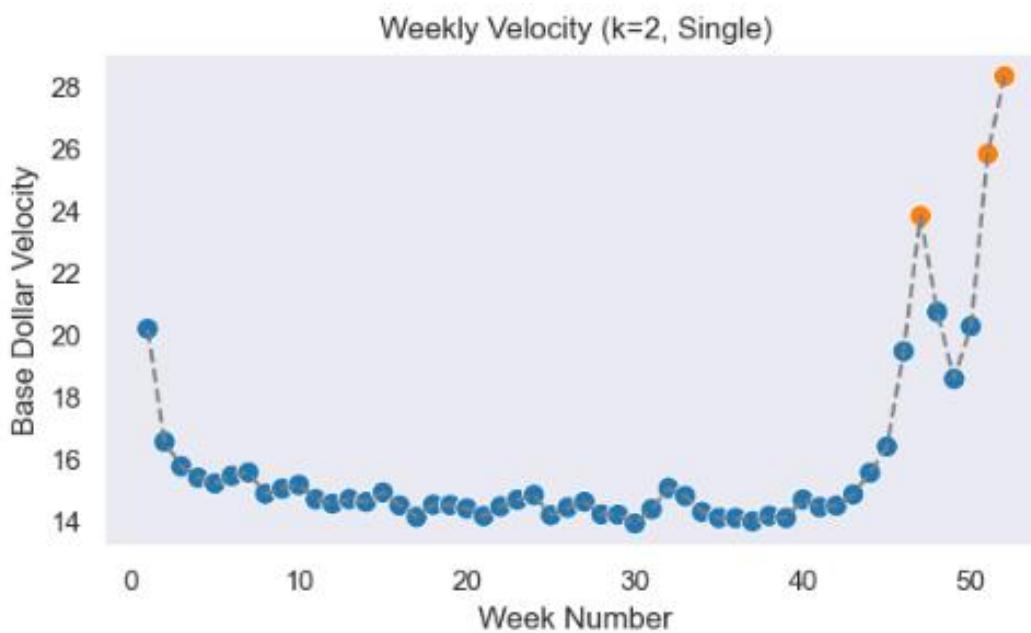


Assessing Results with Pricing Regressions

Case Study: Martinelli's

Takeaway: Clustered seasons can help improve regression R-squared

Martinelli's – Complete linkage appears to cluster seasons more effectively in this case



Martinelli's – Clustered periods yield highest R-squared in pricing regression

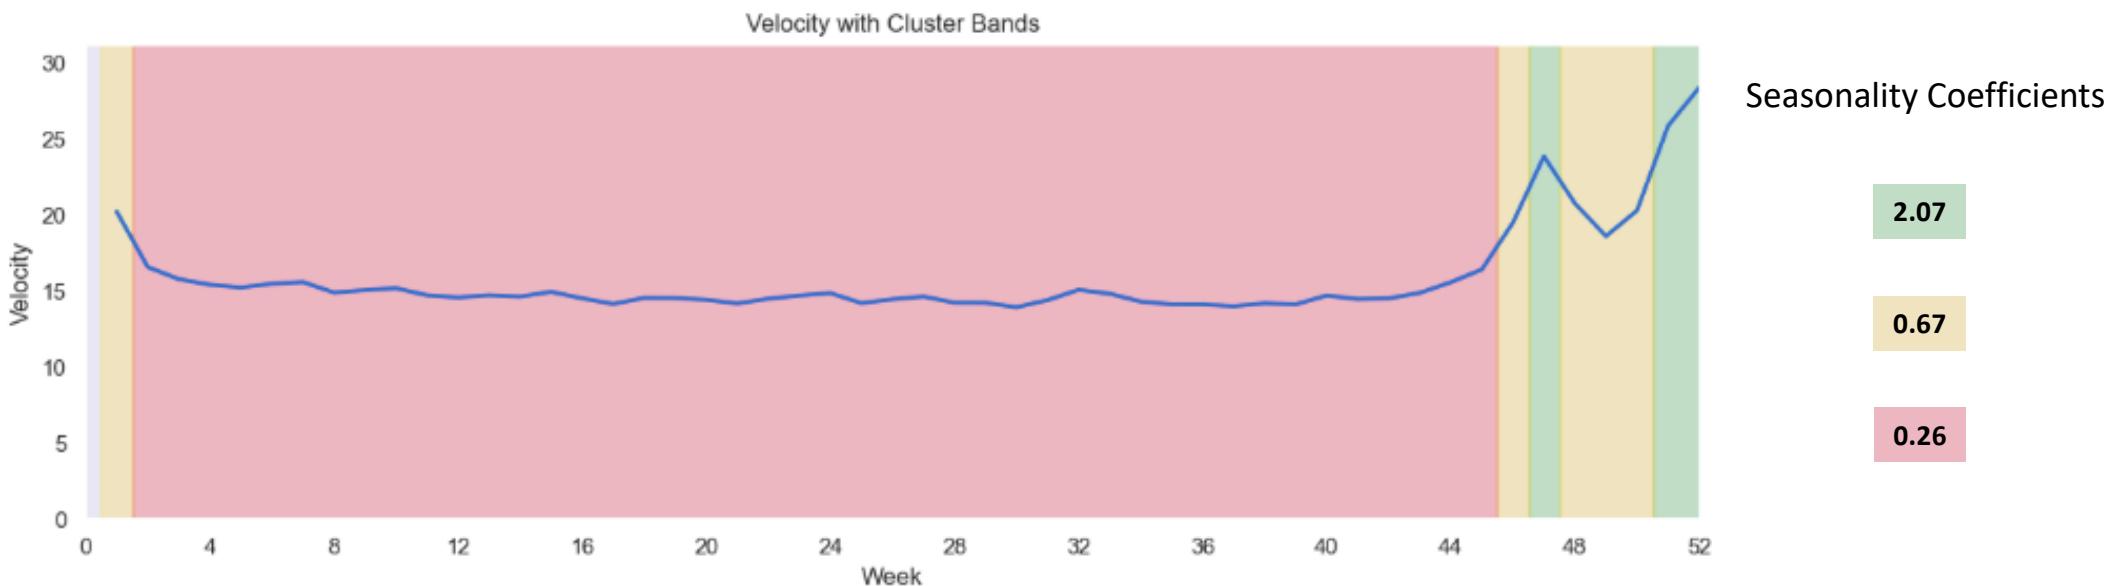
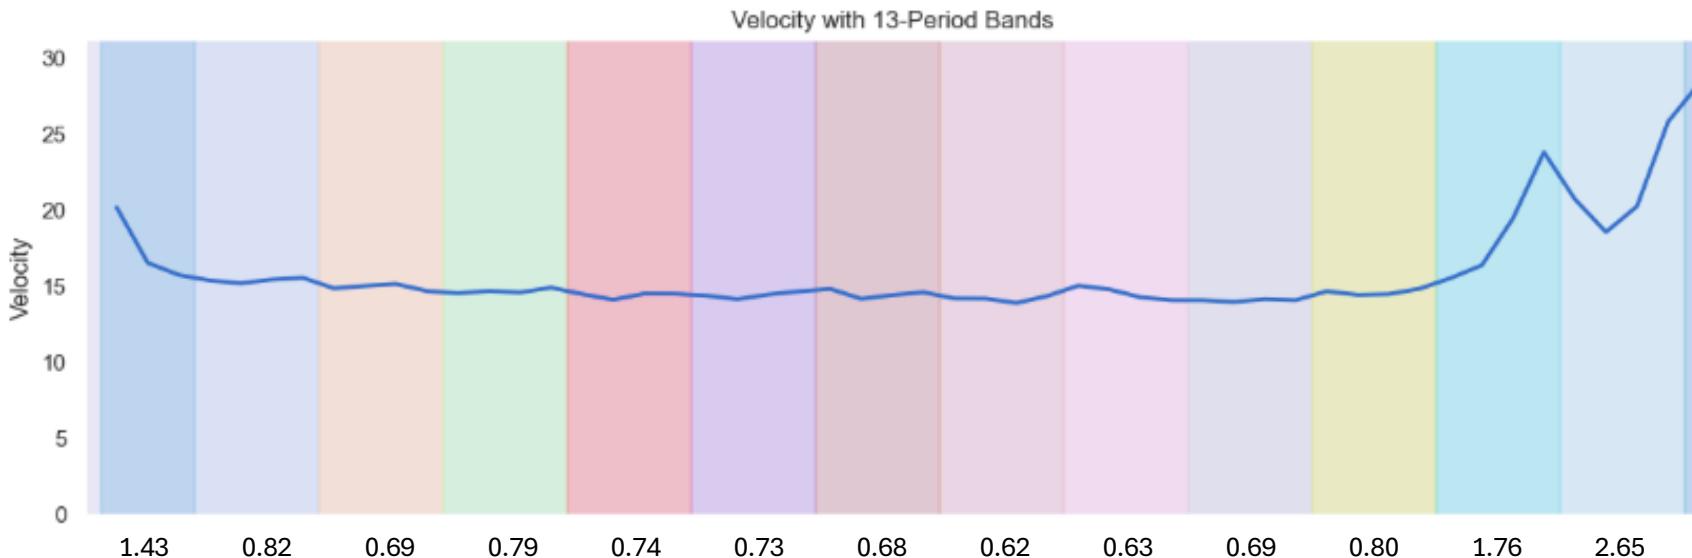
No Period Factor: $\text{np.log}(Q(\text{'Base Units'})) \sim \text{PriceFactor} + \text{np.log}(Q(\text{'ACV'})) + \text{AccountFactor}$

13-Periods Factor: $\text{np.log}(Q(\text{'Base Units'})) \sim \text{PriceFactor} + \text{np.log}(Q(\text{'ACV'})) + \text{AccountFactor} + \text{PeriodFactor}$

3 Clusters Period Factor: $\text{np.log}(Q(\text{'Base Units'})) \sim \text{PriceFactor} + \text{np.log}(Q(\text{'ACV'})) + \text{AccountFactor} + \text{ClusteredPeriodFactor}$

Price	Percent of Base Units	Demand Indices		
		No Period Factor	13-Periods Factor	Clustered Periods Factor
\$2.99	16%	1.00	1.00	1.00
\$3.49	49%	0.65	0.73	0.84
\$3.99	20%	0.37	0.56	0.60
\$4.49	7%	0.31	0.51	0.51
\$4.99	6%	0.27	0.45	0.46
\$5.49	3%	0.25	0.43	0.43
Constant Elasticity		-2.51	-1.36	-1.62
R-squared		0.83	0.90	0.94

Regression on single SKU: Martinelli's Gold Medal Apple Sparkling Cider – Glass Bottle, 25.4 oz (1 ct)



Challenges

- Mapping week numbers across multiple years and aligning holidays consistently
- Many parameters that need to be optimized are at play:
 - Number of clusters
 - Linkage criteria
 - Custom distance metric
 - Alpha value (time vs. velocity weight in chosen custom distance metric)
 - Choice of evaluation metric (Silhouette, CH, DB, regression-based)
- Hard to make a definitive conclusion on impact. Benefits of clustering do not appear to be consistent across datasets and metrics

Continuations

Pipeline

- Implement Python code into Excel for consolidated processing-clustering-regression pipeline.

Regressions

- Determine optimal cluster assignments by maximizing adjusted R-squared or minimizing collinearity metrics (e.g., variance inflation factor, condition number).
- Run the regression pipeline on additional datasets to better quantify the impact of weekly clustering on pricing regressions.
- Compare elasticity predictions from the standard model vs. the clustered model using pre-price-increase data and evaluate which one performs better.

Time-Continuous Seasons

- Explore Markov-constrained clustering to enforce that weeks in the same season are contiguous in time.
- Continue developing and testing custom distance formulas that blend velocity difference with week difference.

Extending Beyond Seasonality – DAFI-Gower Clustering (*Liu et al., 2024*)

- Apply a modified Gower distance that balances the influence of different variable types (numeric, categorical, binary) by scaling them to comparable ranges and weighting them by their importance.
 - Potential variables: velocity, product attributes (package size, flavor, category), promo status, ACV, and account/channel type.

Acknowledgements

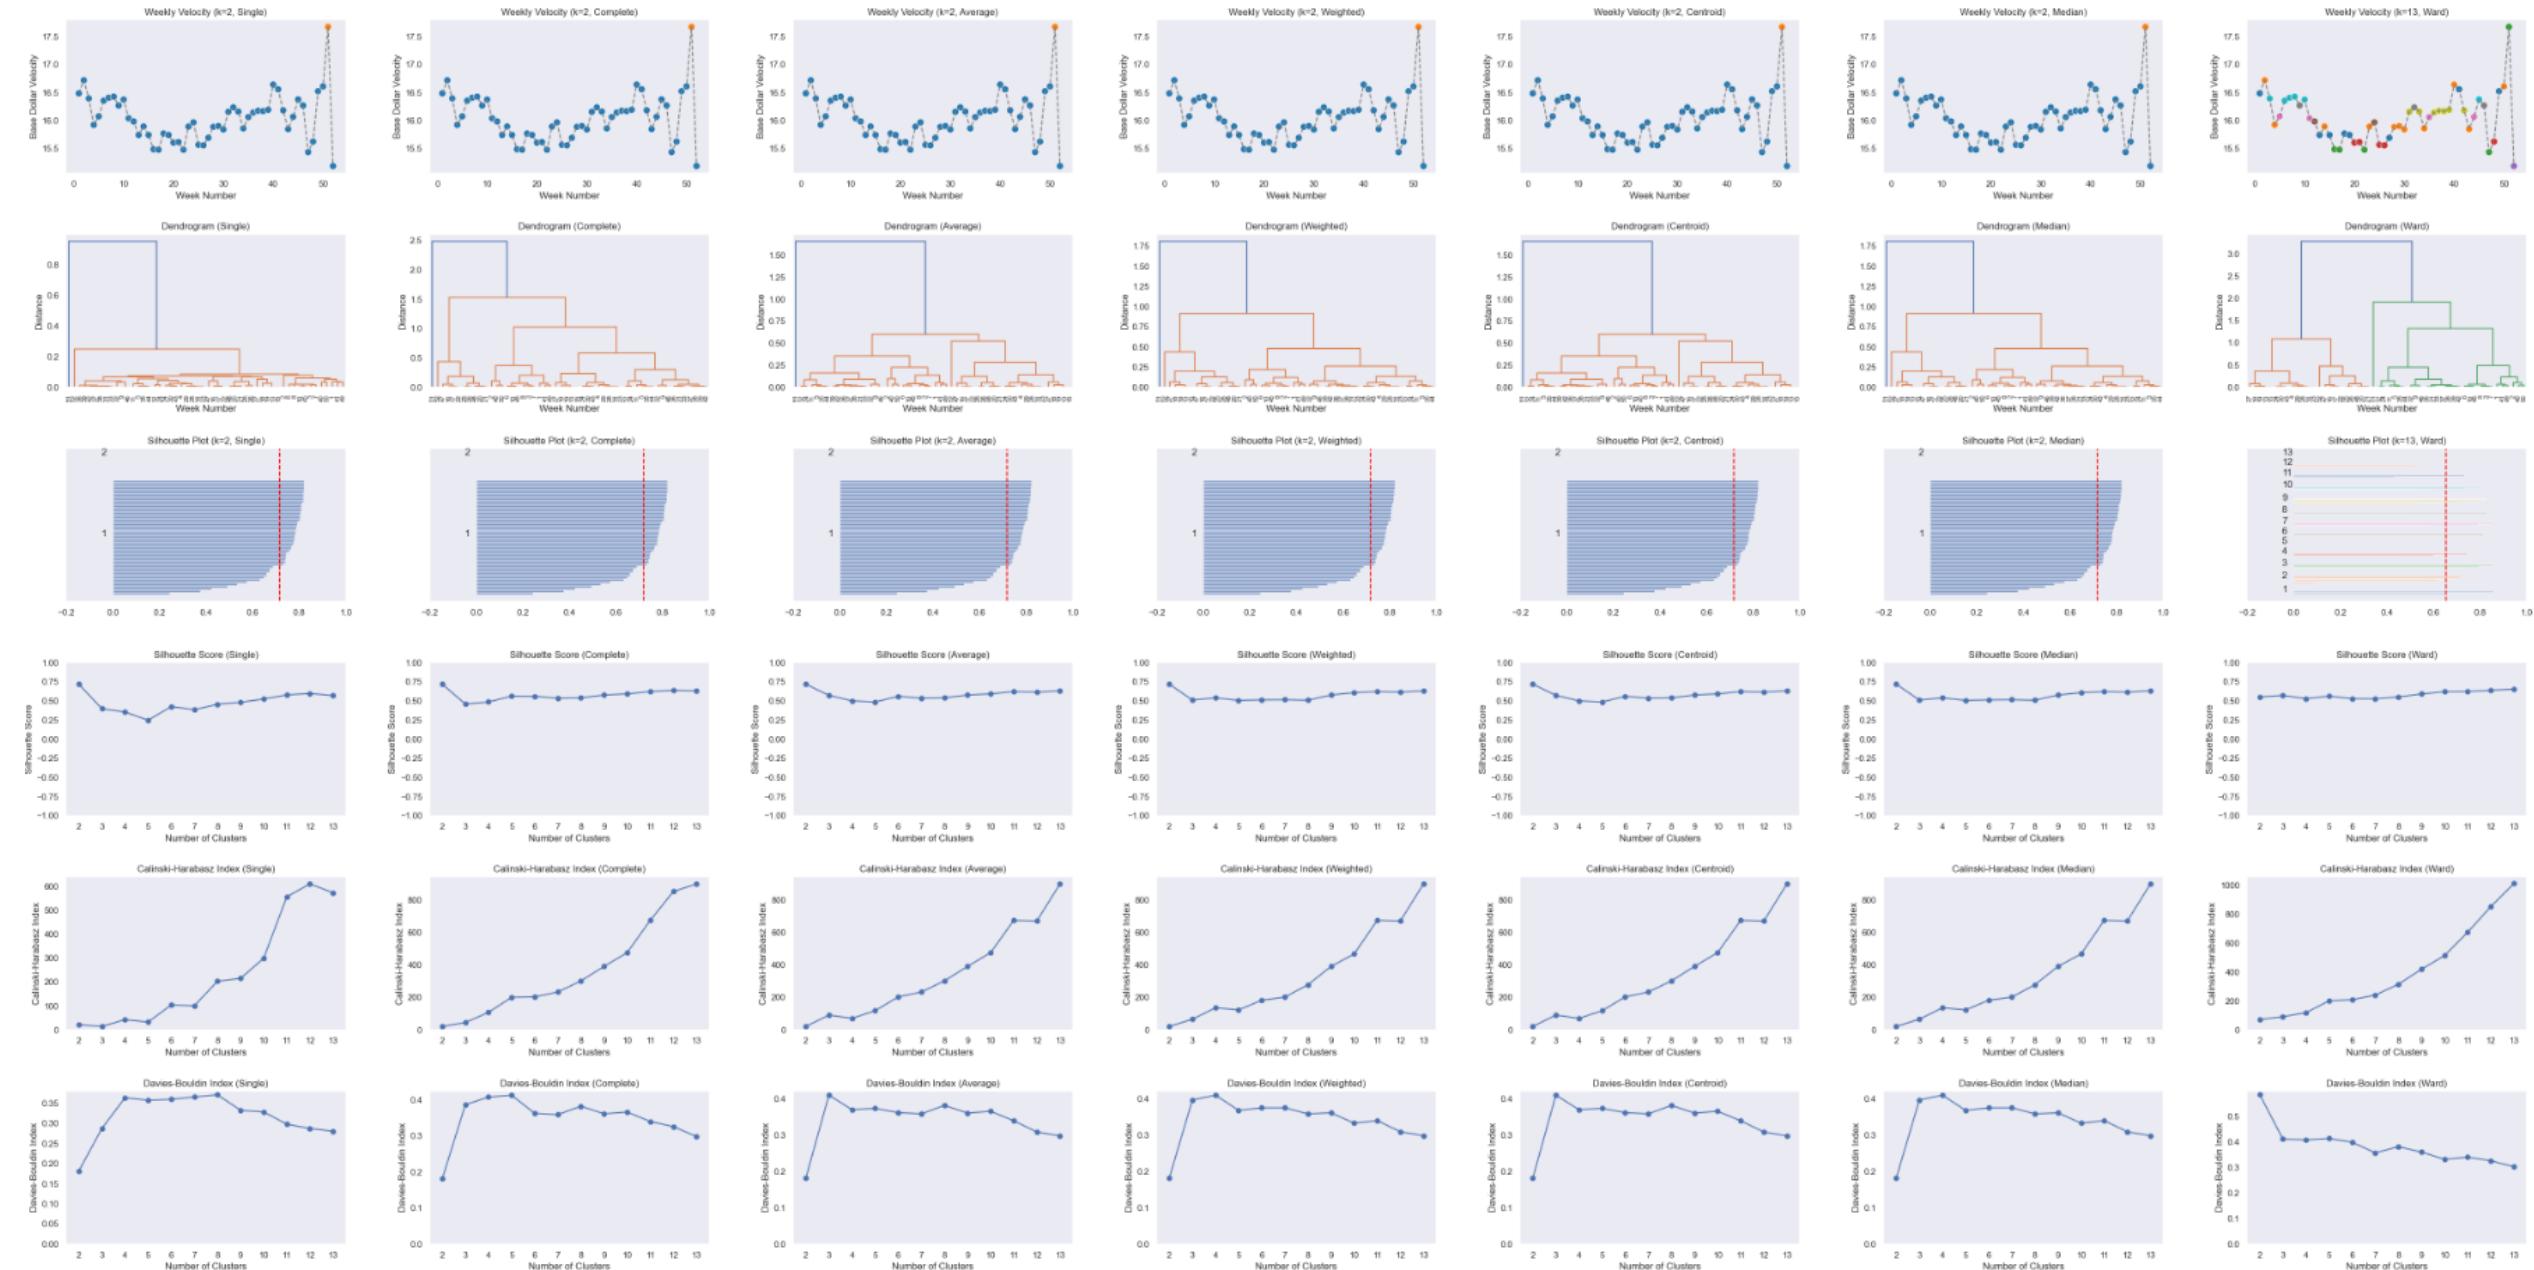
- Sean Dunbar
- William Dumas
- Hazel McCarthy
- Bob Dumas
- Diana Constantinescu

References

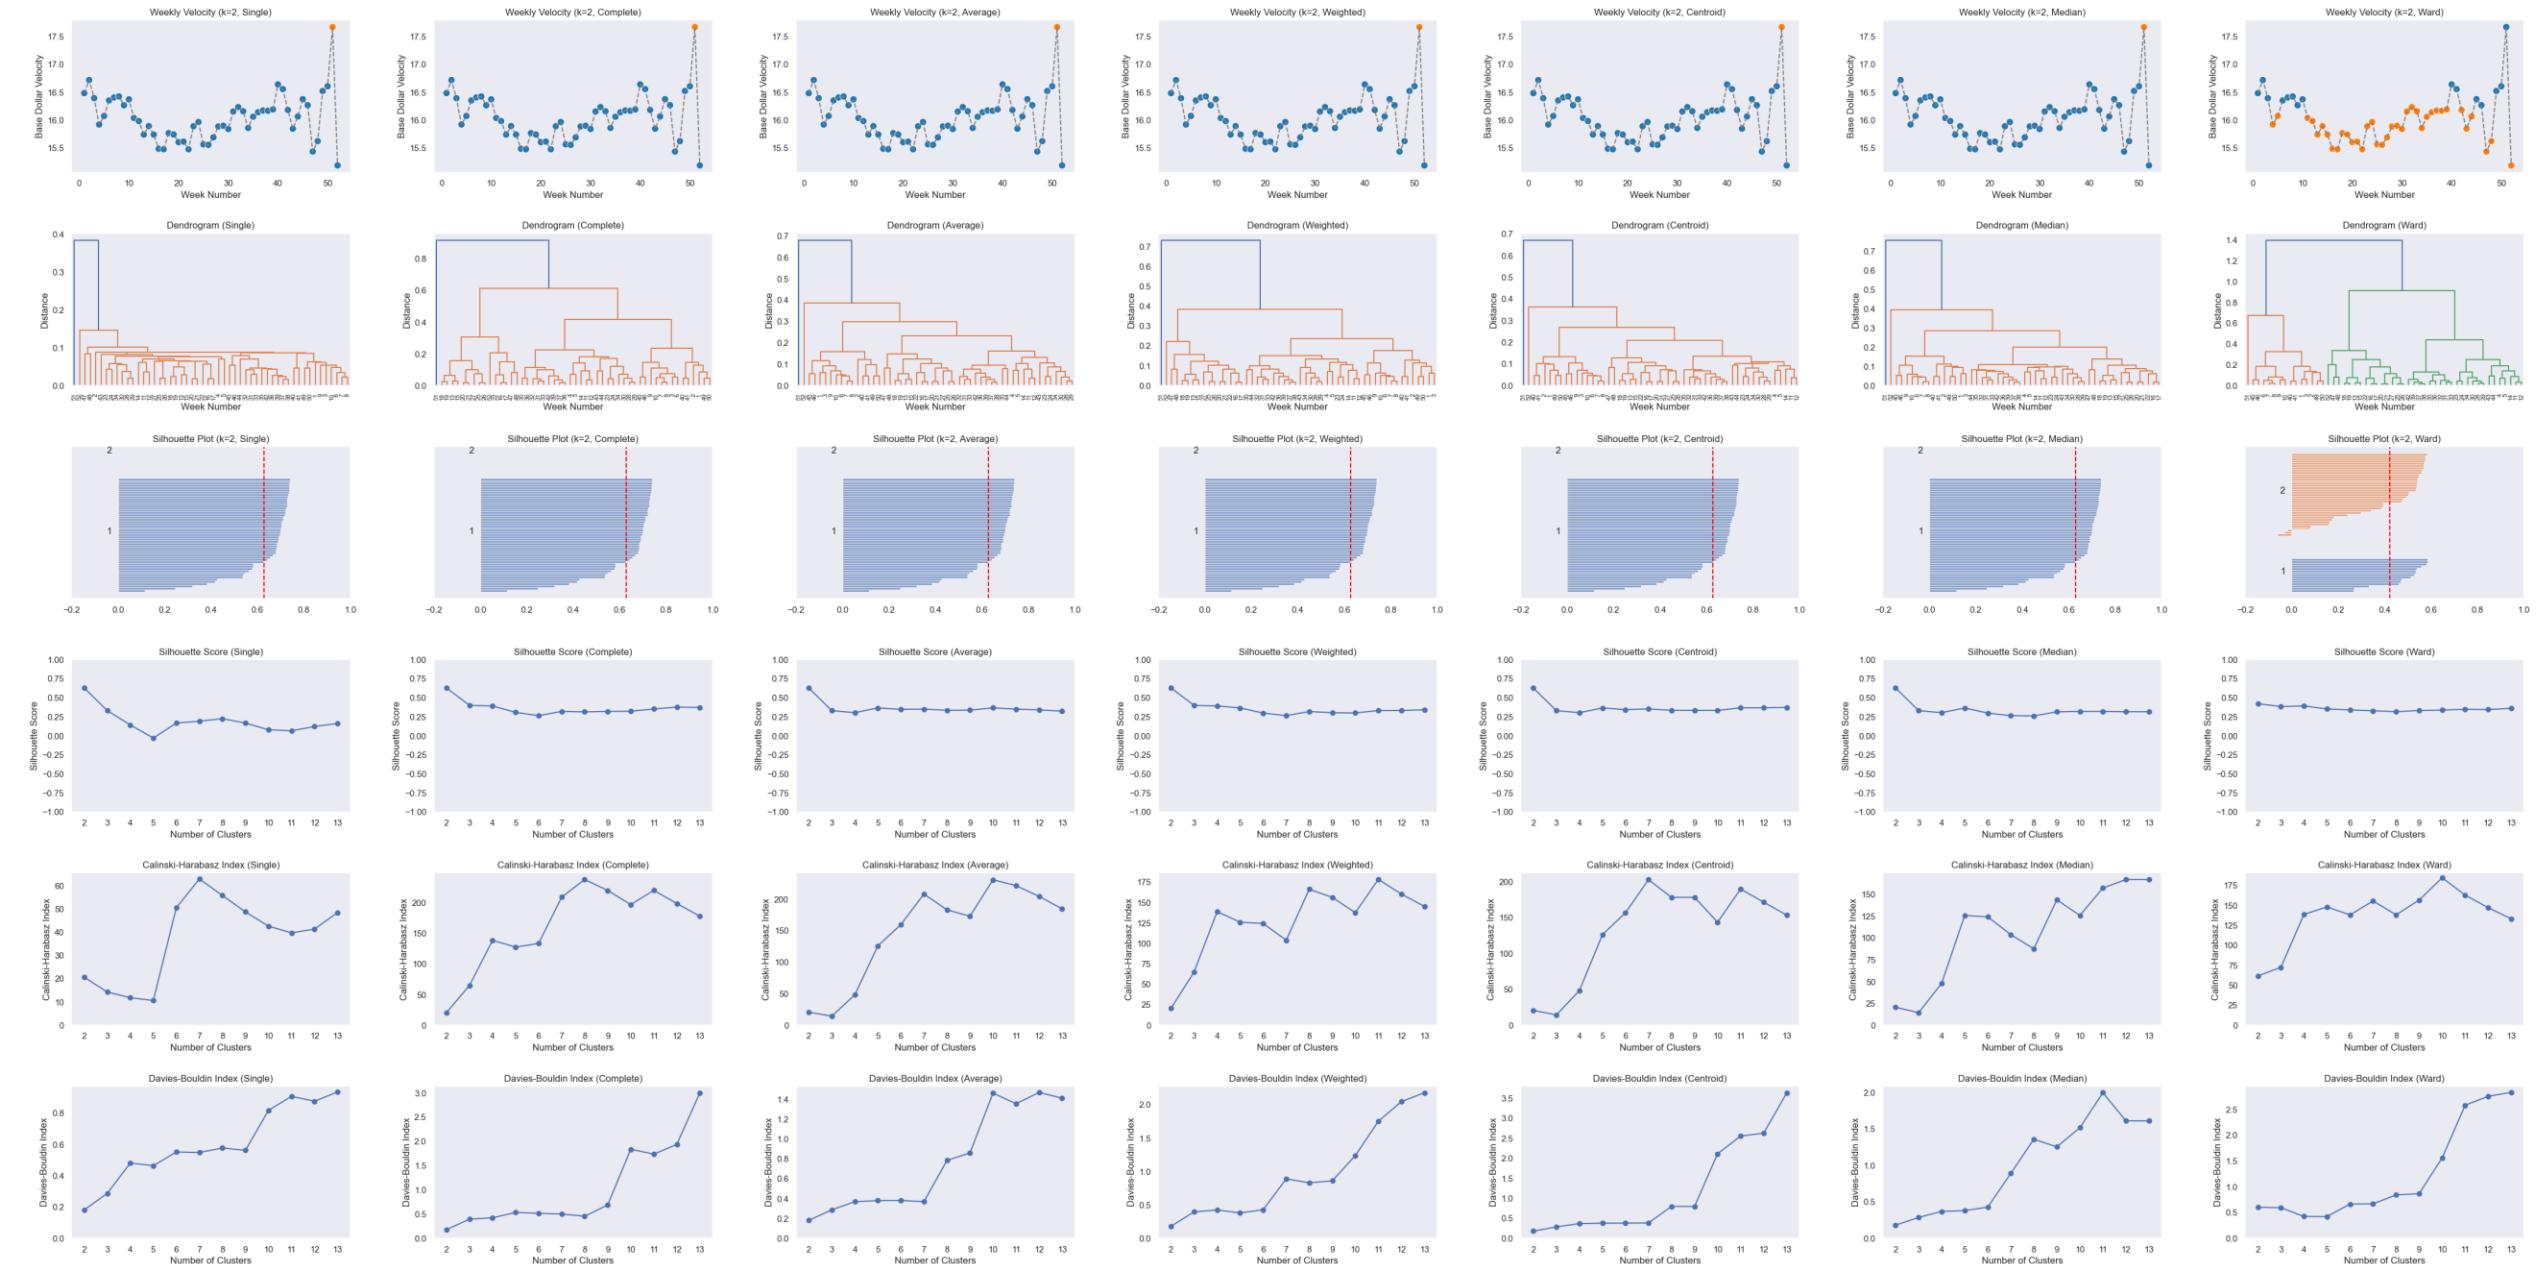
- Liu, P., Yuan, H., Ning, Y. *et al.* A modified and weighted Gower distance-based clustering analysis for mixed type data: a simulation and empirical analyses. *BMC Med Res Methodol* 24, 305 (2024). <https://doi.org/10.1186/s12874-024-02427-8>
- Makkar, A. (2022, December 10). *Vast Data Deduction: Insight on how a Computer Pinpoints Data*. students x students. Retrieved August 12, 2025, from <https://studentsxstudents.com/vast-data-deduction-insight-on-how-a-computer-earmarks-data-af51f9870b58>
- Pedregosa, F., *et al.* (2011). *Scikit-learn: Machine Learning in Python*. Journal of Machine Learning Research, 12, 2825–2830. Retrieved August 12, 2025, from <https://scikit-learn.org/stable/modules/clustering.html>
- Virtanen, P., et al. (2025). *scipy.cluster.hierarchy.linkage* (v1.16.1) [Documentation]. SciPy. Retrieved August 12, 2025, from <https://docs.scipy.org/doc/scipy/reference/generated/scipy.cluster.hierarchy.linkage.html>

Appendix

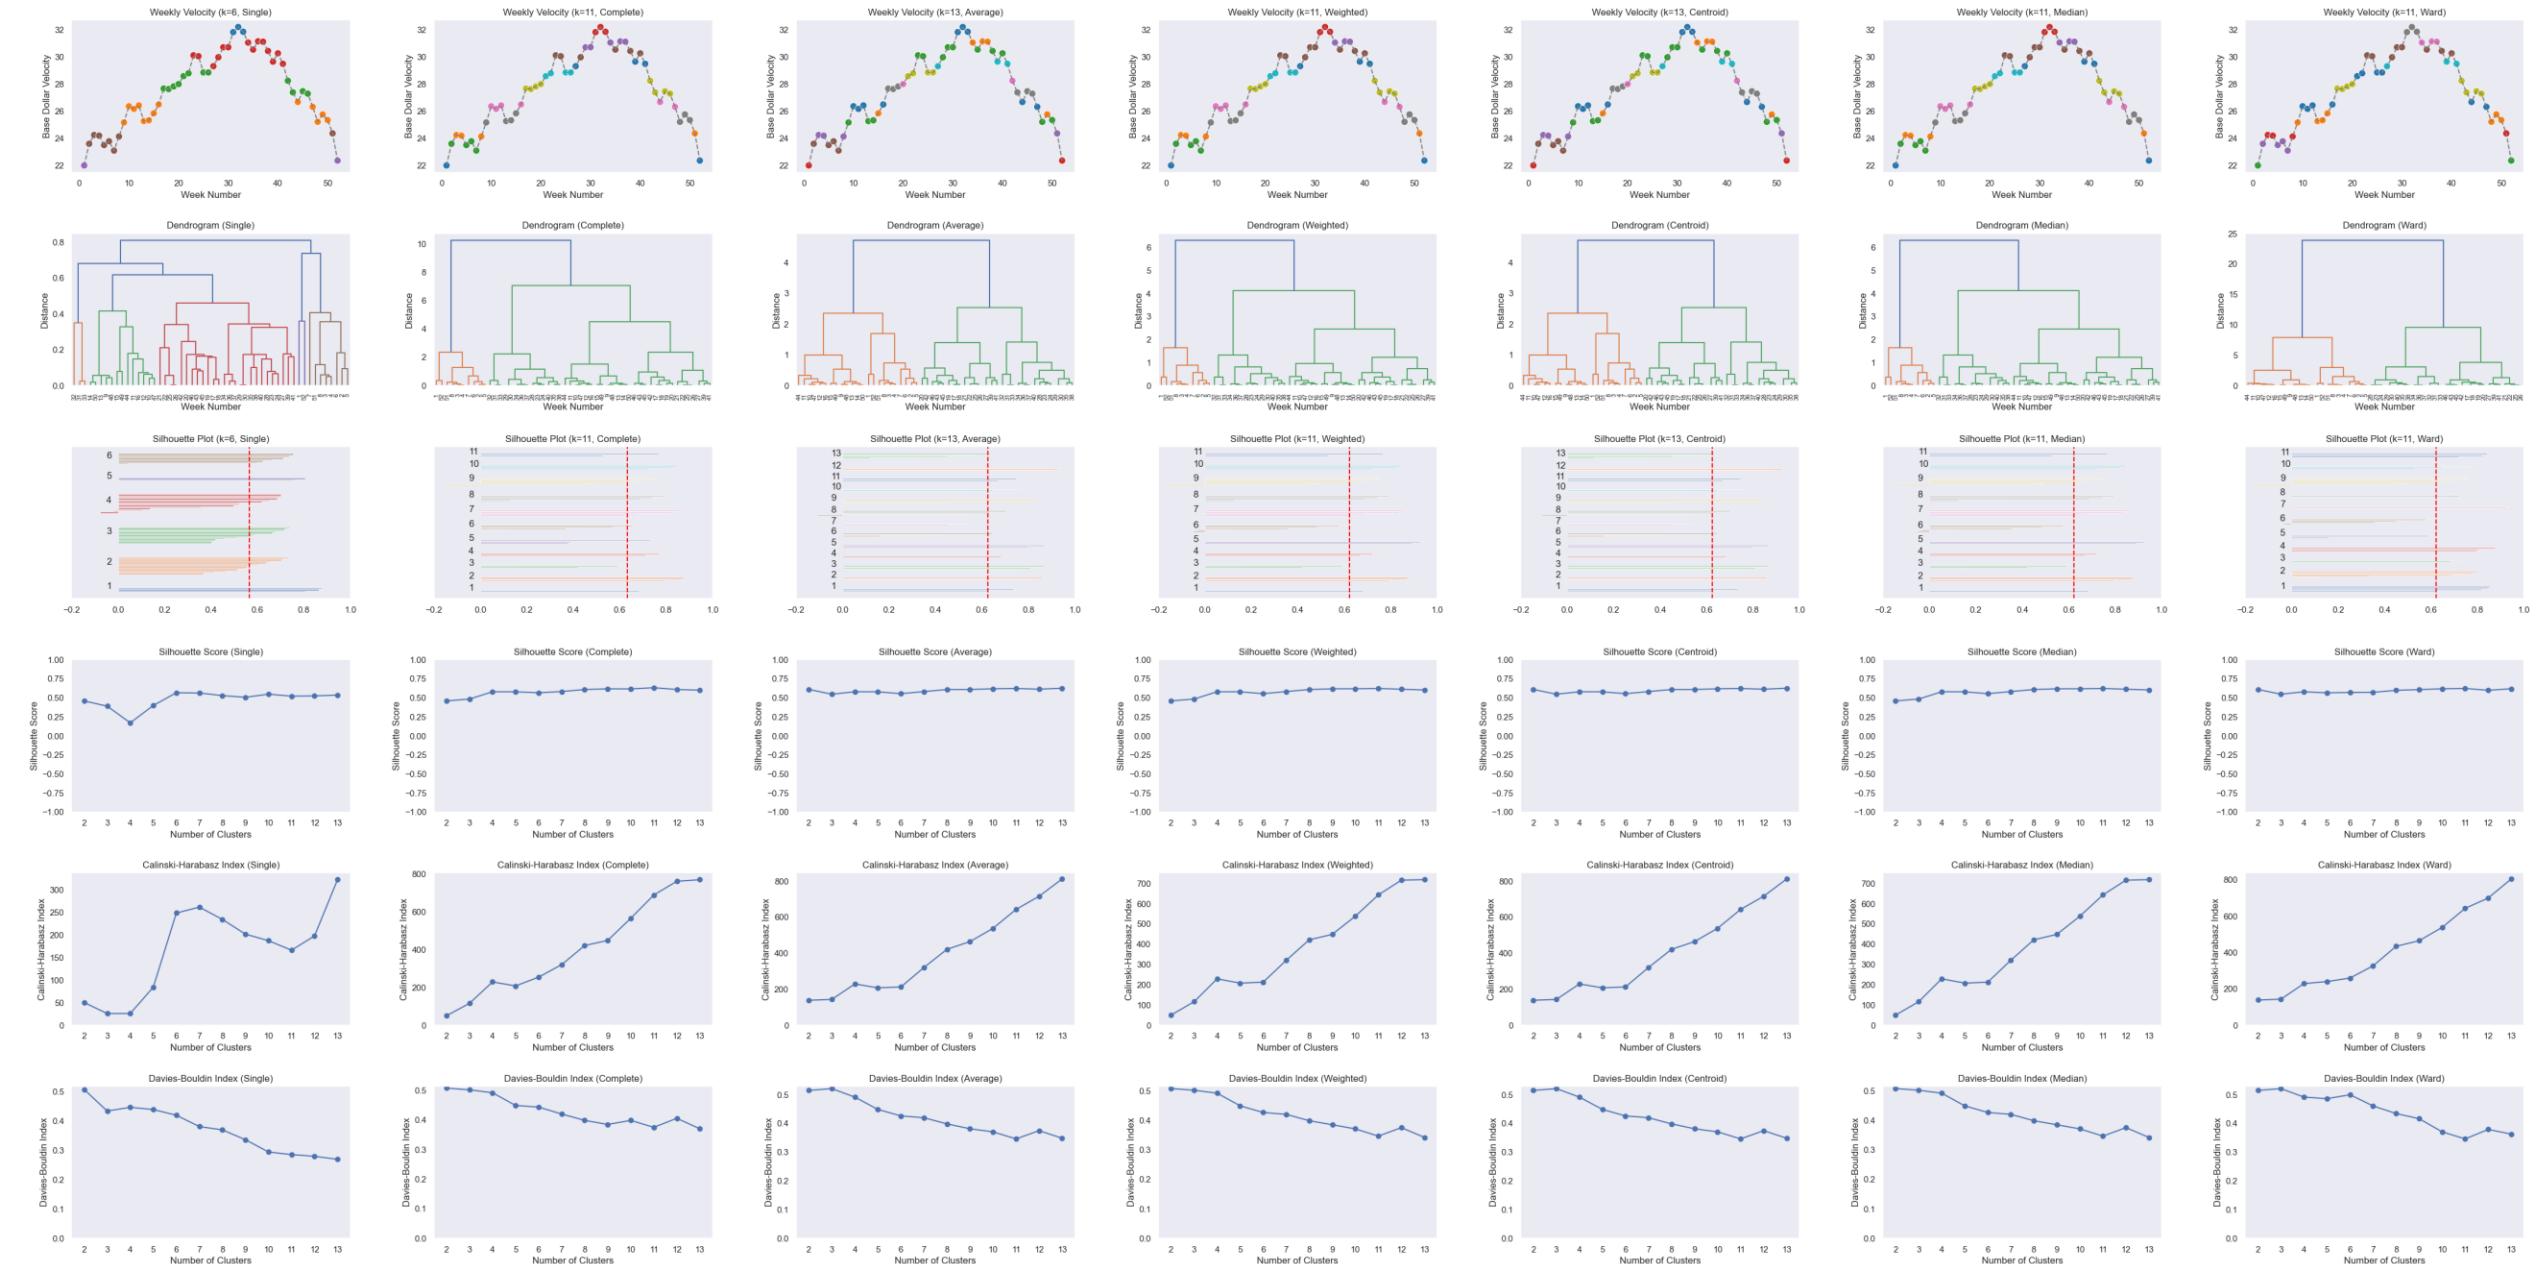
Chomps – Clustering on Velocity



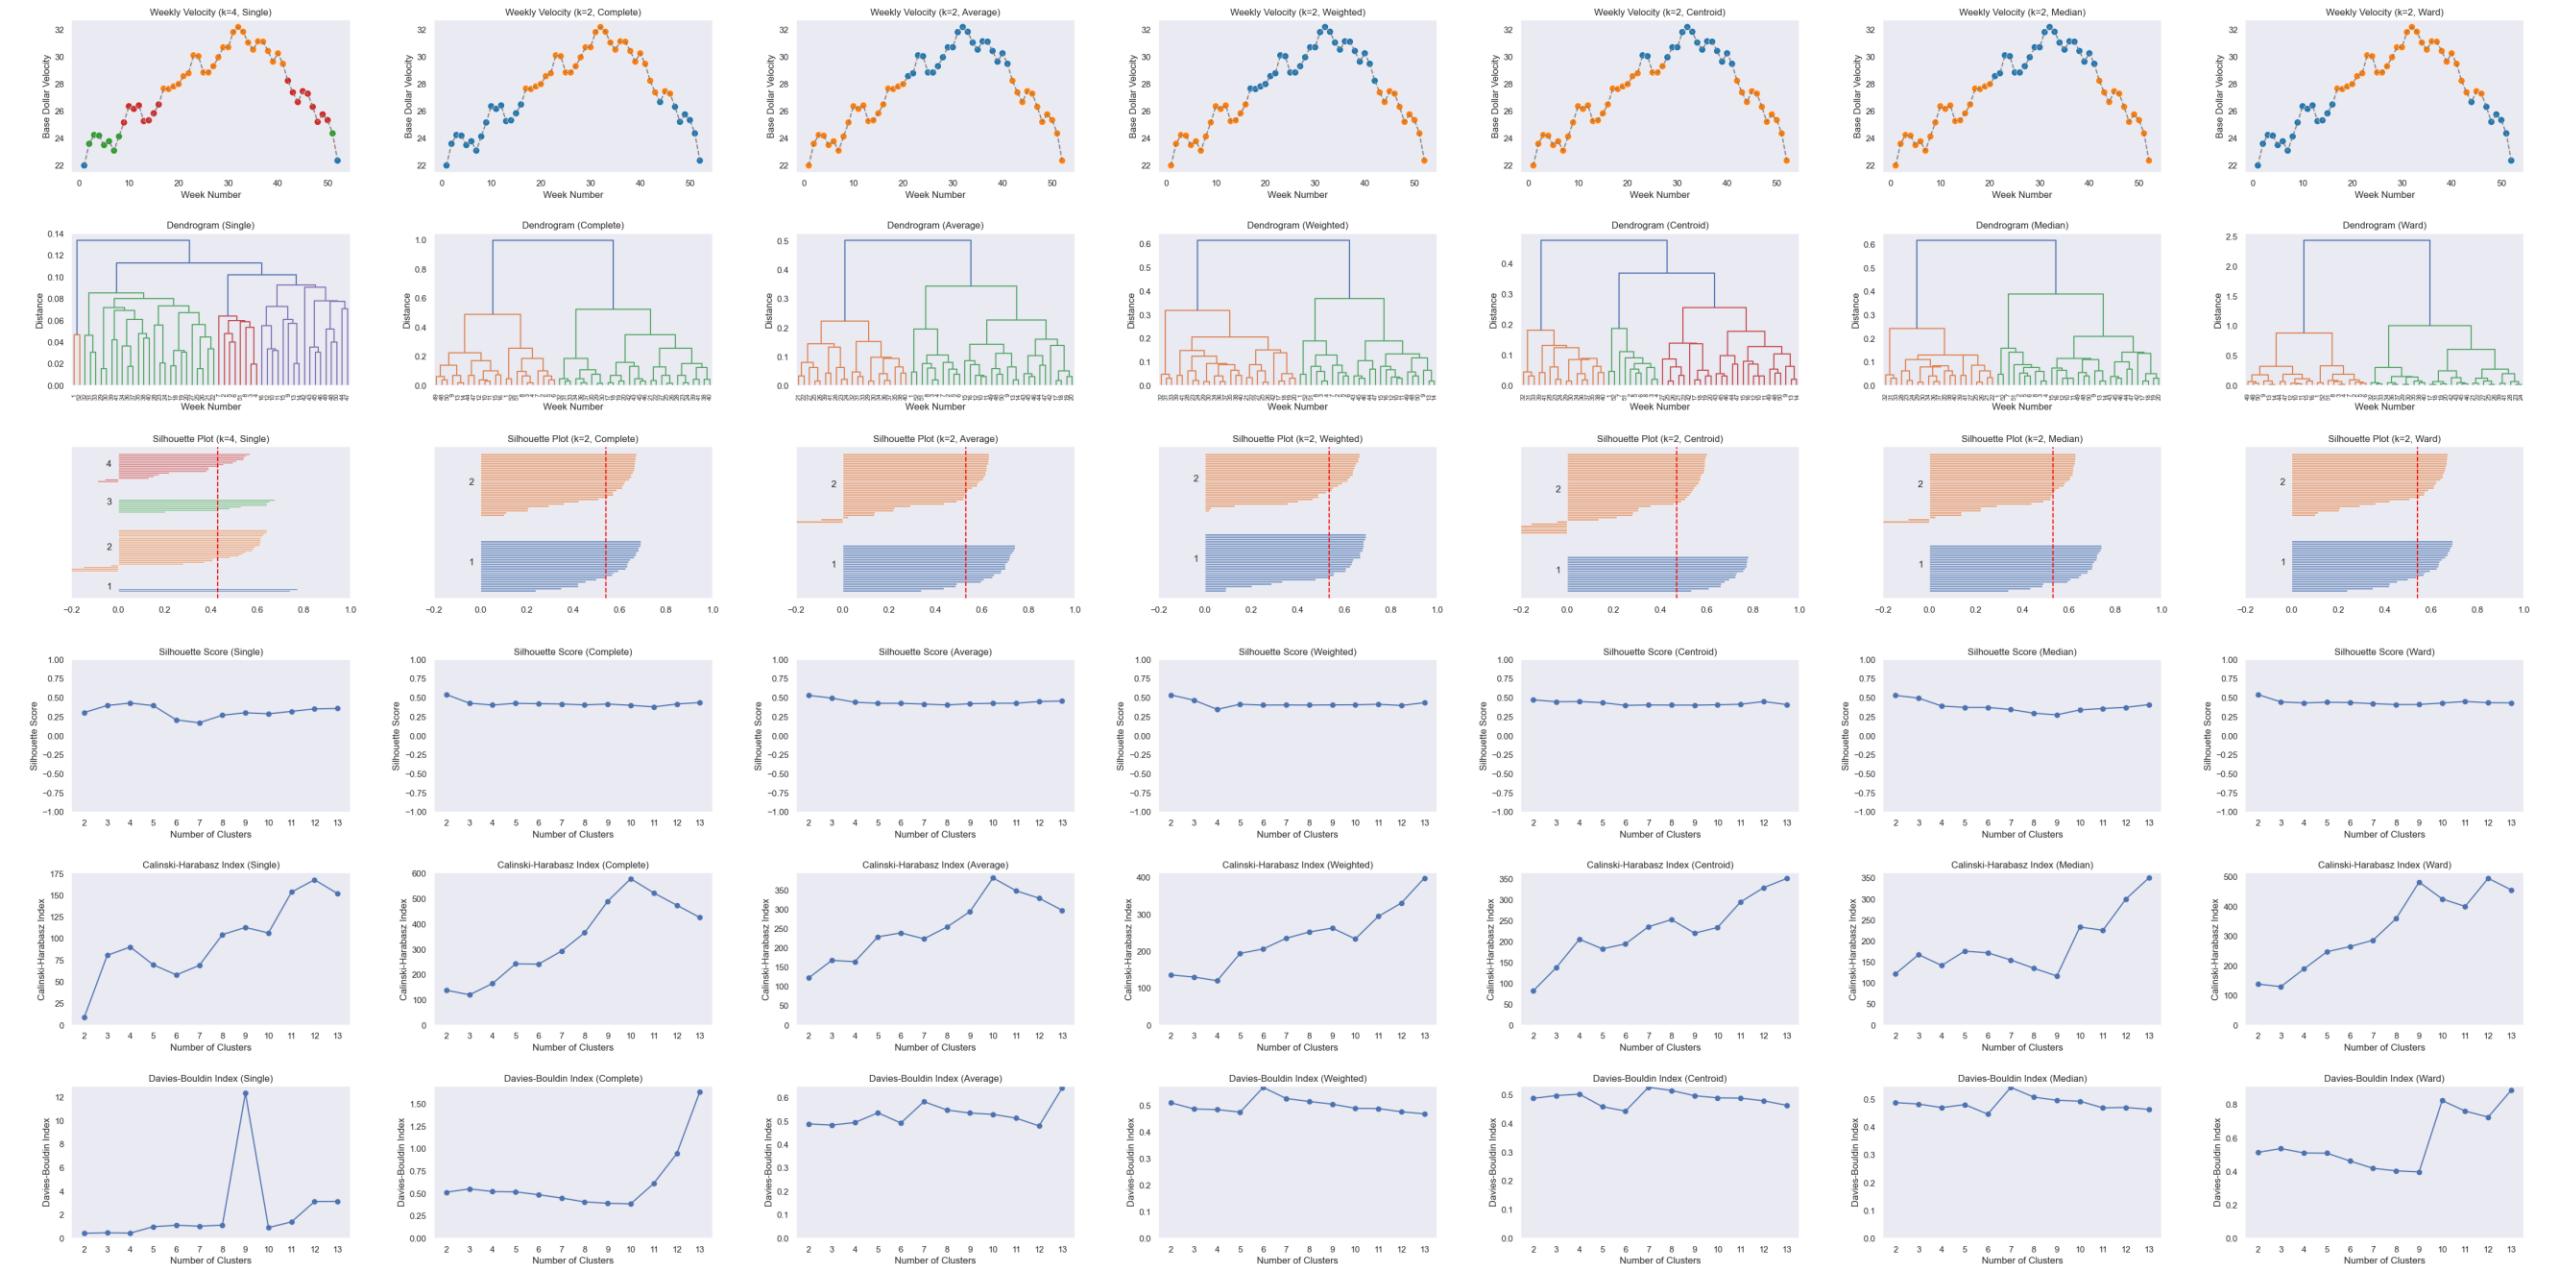
Chomps – Clustering on Velocity & Time (alpha = 0.1, exp. decay distance function)



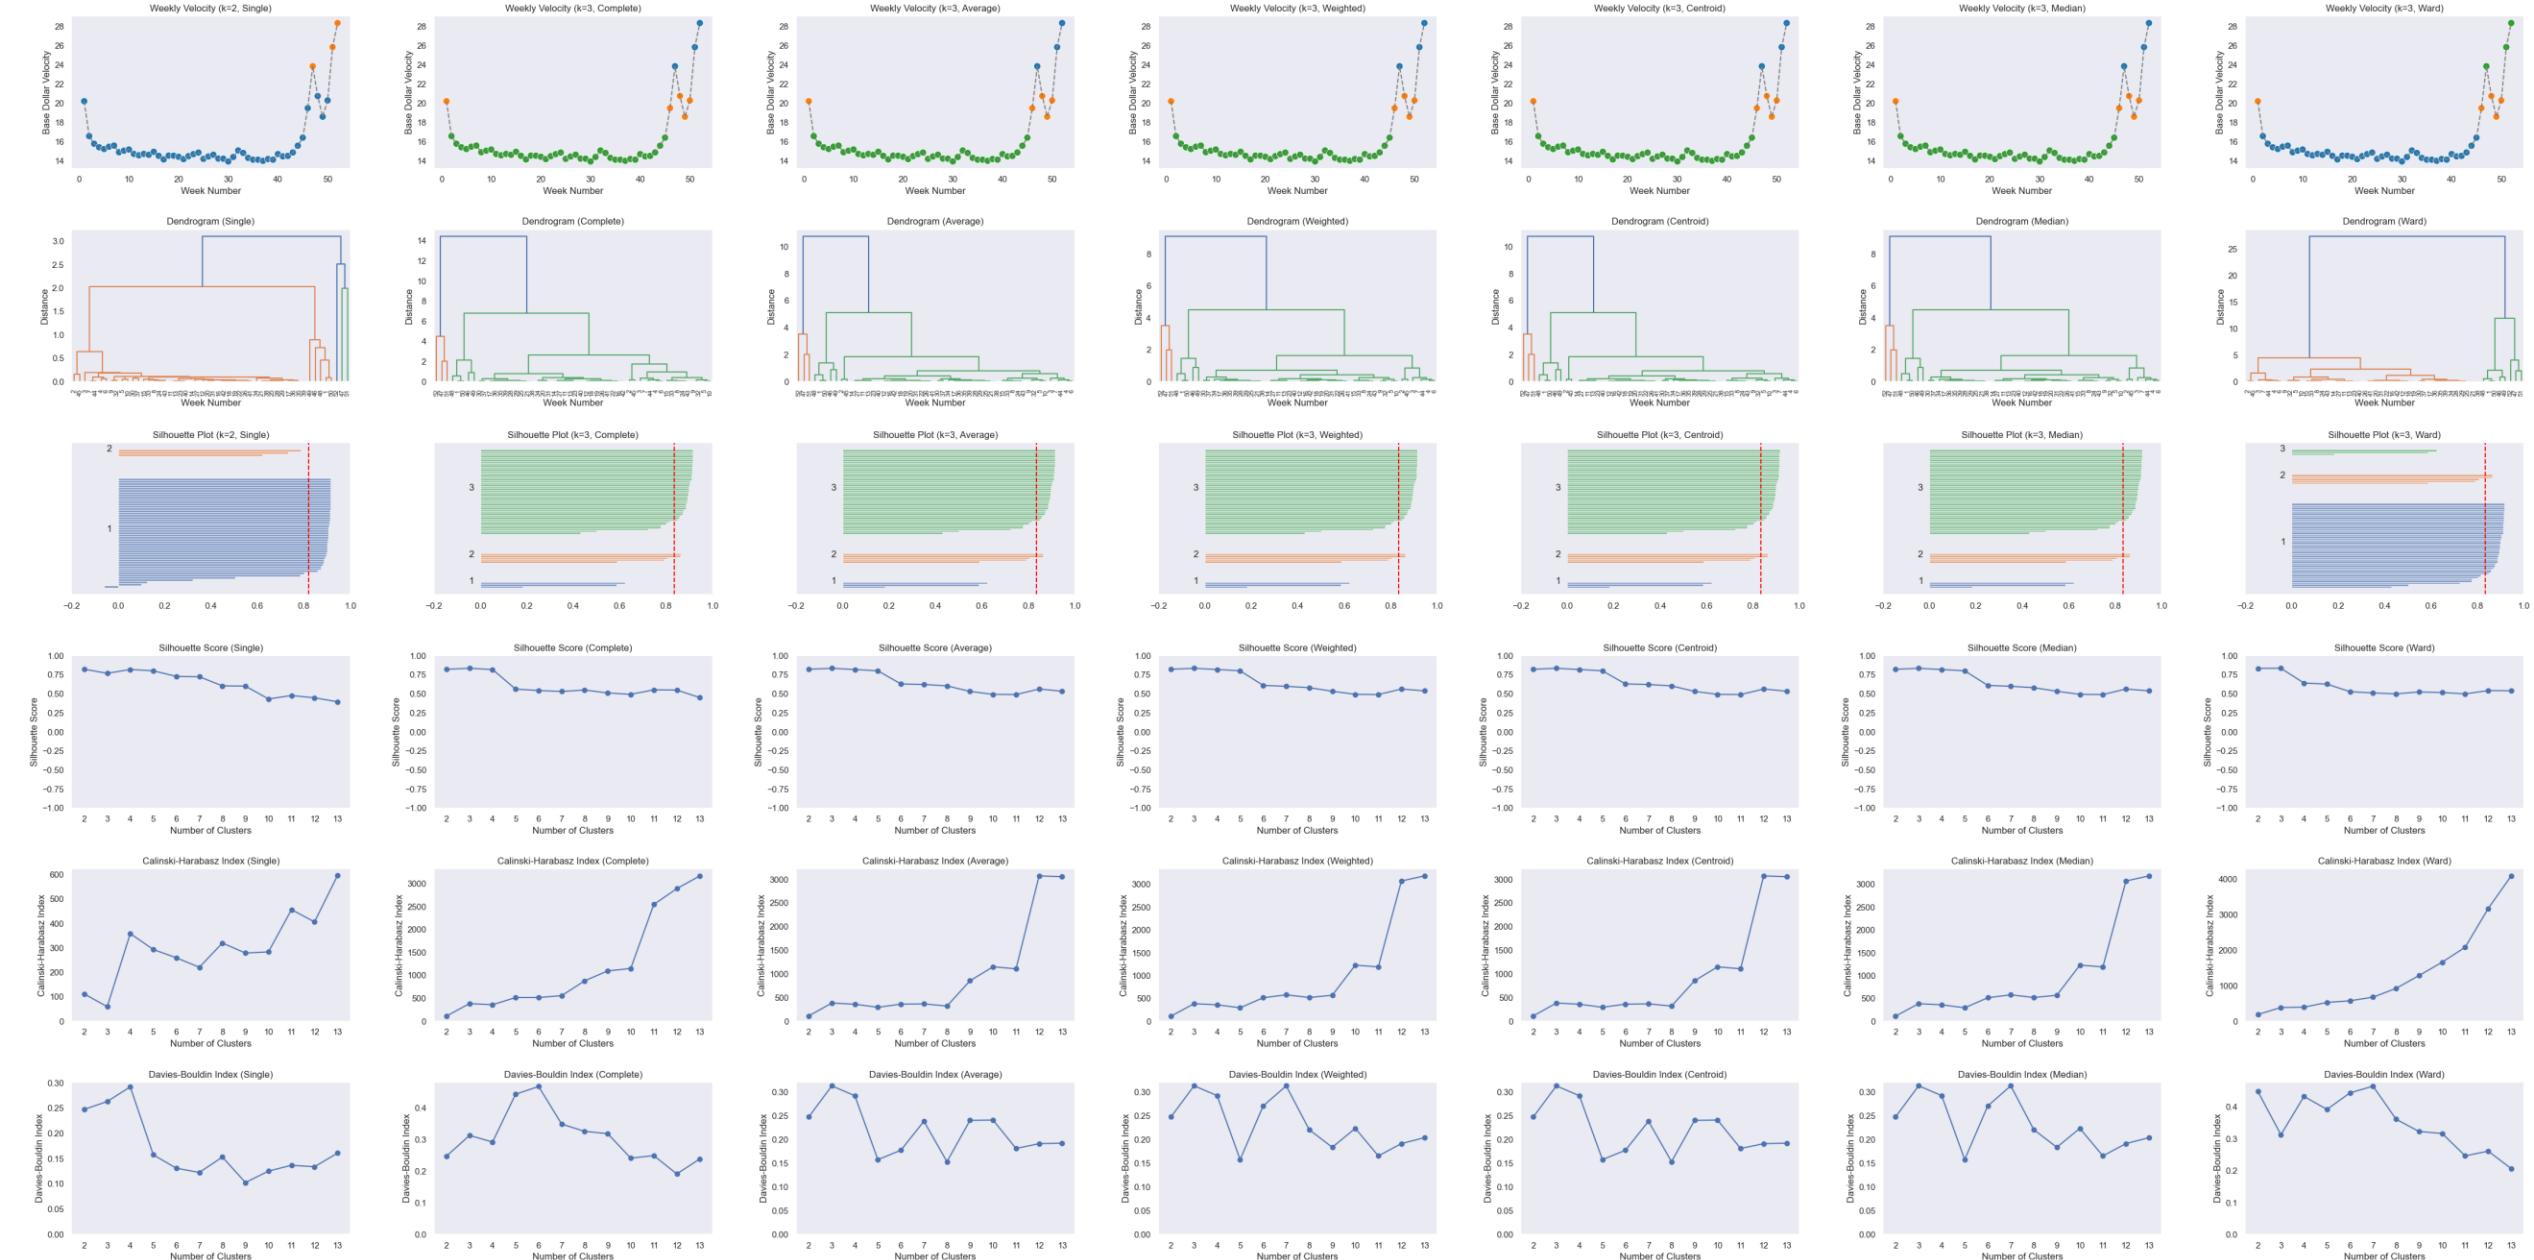
TruFru – Clustering on Velocity



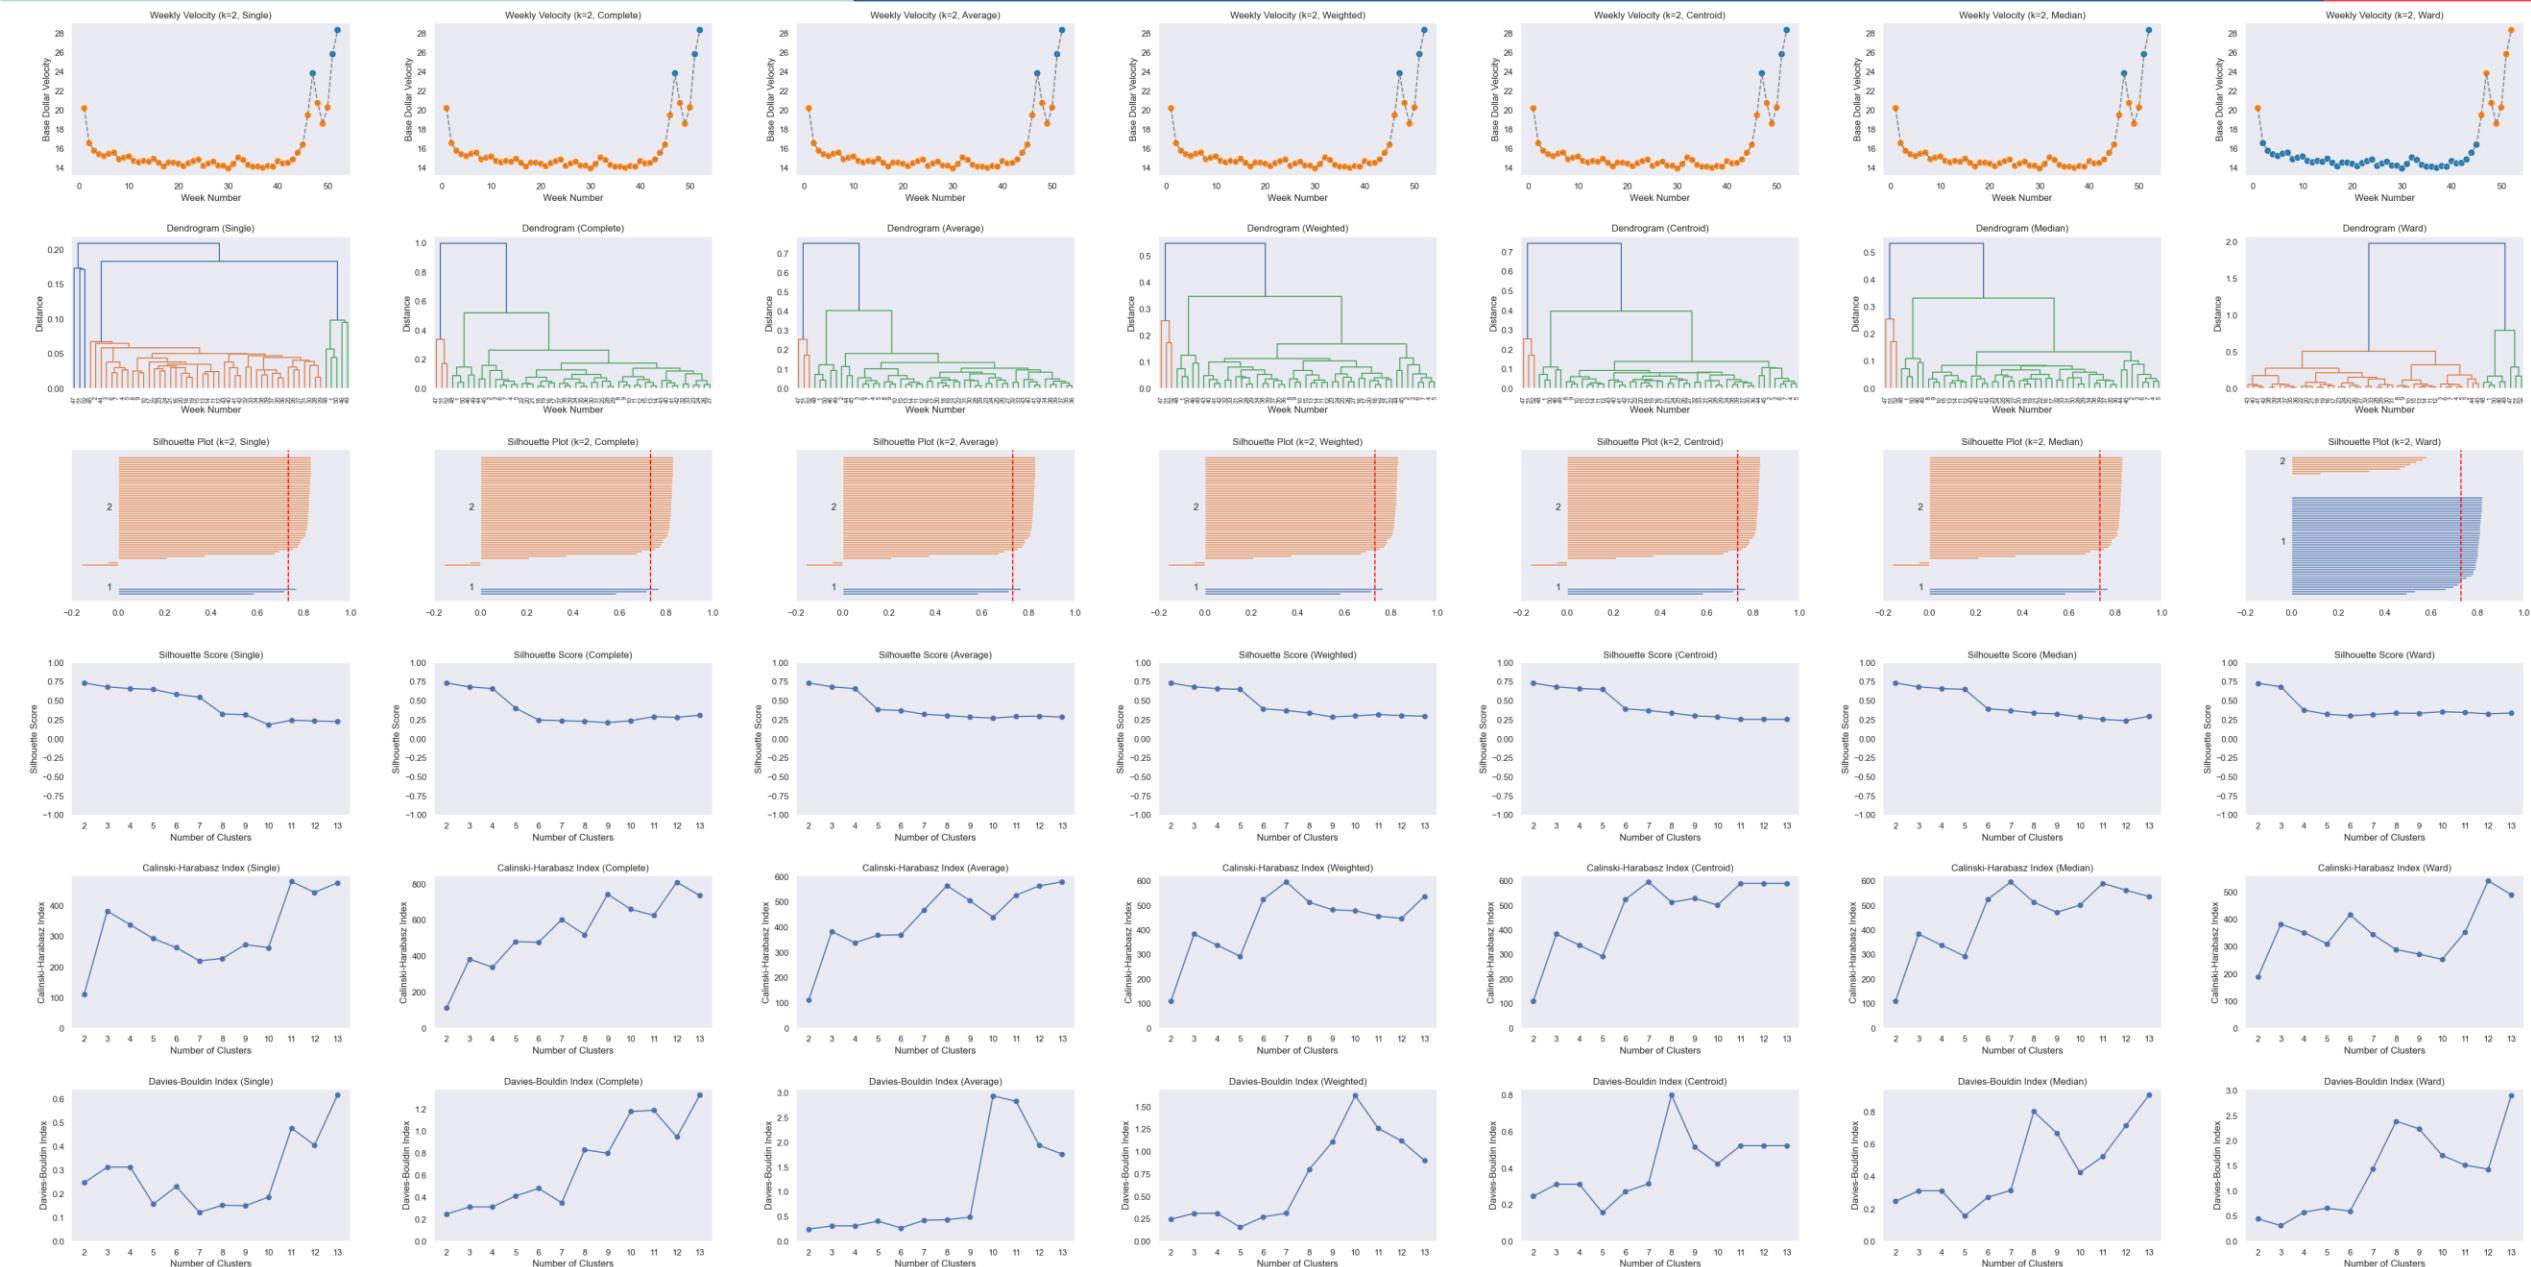
TruFru – Clustering on Velocity & Time (alpha = 0.1, exp. decay distance function)



Martinelli's – Clustering on Velocity



Martinelli's – Clustering on Velocity & Time (alpha = 0.1, exp. decay distance function)



Linkage Criterion

- **Single linkage** – Distance between two clusters = shortest distance between any two points (nearest neighbor). Can create “chained” clusters.
- **Complete linkage** – Distance = farthest distance between any two points (furthest neighbor). Produces compact, evenly shaped clusters.
- **Average linkage** – Distance = average of all pairwise distances between points in the two clusters. Balances chaining and compactness.
- **Weighted linkage** – Like average linkage but updates distances with equal weight to each existing cluster, regardless of size.
- **Centroid linkage** – Distance between clusters = distance between their centroids (mean vectors). Can cause reversals in dendograms.
- **Median linkage** – Similar to centroid but uses median instead of mean for each dimension; more robust to outliers.
- **Ward’s linkage** – Merges clusters that result in the smallest increase in total within-cluster variance; tends to create clusters of similar size.

Linkage Criterion – SciPy Documentation

- method='single' assigns

$$d(u, v) = \min(\text{dist}(u[i], v[j]))$$

for all points i in cluster u and j in cluster v . This is also known as the Nearest Point Algorithm.

- method='complete' assigns

$$d(u, v) = \max(\text{dist}(u[i], v[j]))$$

for all points i in cluster u and j in cluster v . This is also known by the Farthest Point Algorithm or Voor Hees Algorithm.

- method='average' assigns

$$d(u, v) = \sum_{ij} \frac{d(u[i], v[j])}{(|u| * |v|)}$$

for all points i and j where $|u|$ and $|v|$ are the cardinalities of clusters u and v , respectively. This is also called the UPGMA algorithm.

- method='weighted' assigns

$$d(u, v) = (\text{dist}(s, v) + \text{dist}(t, v))/2$$

where cluster u was formed with cluster s and t and v is a remaining cluster in the forest (also called WPGMA).

- method='centroid' assigns

$$\text{dist}(s, t) = \|c_s - c_t\|_2$$

where c_s and c_t are the centroids of clusters s and t , respectively. When two clusters s and t are combined into a new cluster u , the new centroid is computed over all the original objects in clusters s and t . The distance then becomes the Euclidean distance between the centroid of u and the centroid of a remaining cluster v in the forest. This is also known as the UPGMC algorithm.

- method='median' assigns $d(s, t)$ like the **centroid** method. When two clusters s and t are combined into a new cluster u , the average of centroids s and t give the new centroid u . This is also known as the WPGMC algorithm.
- method='ward' uses the Ward variance minimization algorithm. The new entry $d(u, v)$ is computed as follows,

$$d(u, v) = \sqrt{\frac{|v| + |s|}{T} d(v, s)^2 + \frac{|v| + |t|}{T} d(v, t)^2 - \frac{|v|}{T} d(s, t)^2}$$

where u is the newly joined cluster consisting of clusters s and t , v is an unused cluster in the forest, $T = |v| + |s| + |t|$, and $| * |$ is the cardinality of its argument. This is also known as the incremental algorithm.

Cluster Evaluation Metrics – scikit-learn Documentation

Silhouette Score

- **a:** The mean distance between a sample and all other points in the same class.
- **b:** The mean distance between a sample and all other points in the *next nearest cluster*.

The Silhouette Coefficient s for a single sample is then given as:

$$s = \frac{b - a}{\max(a, b)}$$

The Silhouette Coefficient for a set of samples is given as the mean of the Silhouette Coefficient for each sample.

Cluster Evaluation Metrics – scikit-learn Documentation

CH Index

For a set of data E of size n_E which has been clustered into k clusters, the Calinski-Harabasz score s is defined as the ratio of the between-clusters dispersion mean and the within-cluster dispersion:

$$s = \frac{\text{tr}(B_k)}{\text{tr}(W_k)} \times \frac{n_E - k}{k - 1}$$

where $\text{tr}(B_k)$ is trace of the between group dispersion matrix and $\text{tr}(W_k)$ is the trace of the within-cluster dispersion matrix defined by:

$$W_k = \sum_{q=1}^k \sum_{x \in C_q} (x - c_q)(x - c_q)^T$$

$$B_k = \sum_{q=1}^k n_q (c_q - c_E)(c_q - c_E)^T$$

with C_q the set of points in cluster q , c_q the center of cluster q , c_E the center of E , and n_q the number of points in cluster q .

Cluster Evaluation Metrics – scikit-learn Documentation

DB Index

The index is defined as the average similarity between each cluster C_i for $i = 1, \dots, k$ and its most similar one C_j . In the context of this index, similarity is defined as a measure R_{ij} that trades off:

- s_i , the average distance between each point of cluster i and the centroid of that cluster – also known as cluster diameter.
- d_{ij} , the distance between cluster centroids i and j .

A simple choice to construct R_{ij} so that it is nonnegative and symmetric is:

$$R_{ij} = \frac{s_i + s_j}{d_{ij}}$$

Then the Davies-Bouldin index is defined as:

$$DB = \frac{1}{k} \sum_{i=1}^k \max_{i \neq j} R_{ij}$$

Custom Distance Functions

```
def dist_linear(obs_1, obs_2, alpha, min_velocity, max_velocity):
    # linear distances
    week_diff = min(abs(obs_1[0] - obs_2[0]), 52 - abs(obs_1[0] - obs_2[0]))
    week_norm = week_diff / 26
    vel_1 = obs_1[1]
    vel_2 = obs_2[1]
    vel_norm = abs(vel_1 - vel_2) / (max_velocity - min_velocity + 1e-8)
    return alpha * week_norm + (1 - alpha) * vel_norm

def dist_squared(obs_1, obs_2, alpha, min_velocity, max_velocity):
    # squared velocity distance
    week_diff = min(abs(obs_1[0] - obs_2[0]), 52 - abs(obs_1[0] - obs_2[0]))
    week_norm = week_diff / 26
    vel_1 = obs_1[1]
    vel_2 = obs_2[1]
    vel_norm = abs(vel_1 - vel_2) / (max_velocity - min_velocity)
    vel_sq = vel_norm ** 2
    return alpha * week_norm + (1-alpha) * vel_sq
```

```
def dist_exp_decay(obs_1, obs_2, alpha, min_velocity, max_velocity):
    # exponential decay on time
    week_diff = min(abs(obs_1[0] - obs_2[0]), 52 - abs(obs_1[0] - obs_2[0]))
    tau = 6
    week_decay = 1 - math.exp(-week_diff / tau)
    vel_1 = obs_1[1]
    vel_2 = obs_2[1]
    vel_norm = abs(vel_1 - vel_2) / (max_velocity - min_velocity + 1e-8)
    return alpha * week_decay + (1 - alpha) * vel_norm

def dist_cosine(obs_1, obs_2, alpha, min_velocity, max_velocity):
    # cosine on time within surrounding 4 weeks, then linear
    week_diff = min(abs(obs_1[0] - obs_2[0]), 52 - abs(obs_1[0] - obs_2[0]))
    width = 4
    if week_diff <= width:
        week_bump = 0.5 * (1 - math.cos(math.pi * week_diff / width))
    else:
        week_bump = week_diff / 26
    vel_1 = obs_1[1]
    vel_2 = obs_2[1]
    vel_norm = abs(vel_1 - vel_2) / (max_velocity - min_velocity + 1e-8)
    return alpha * week_bump + (1 - alpha) * vel_norm
```